

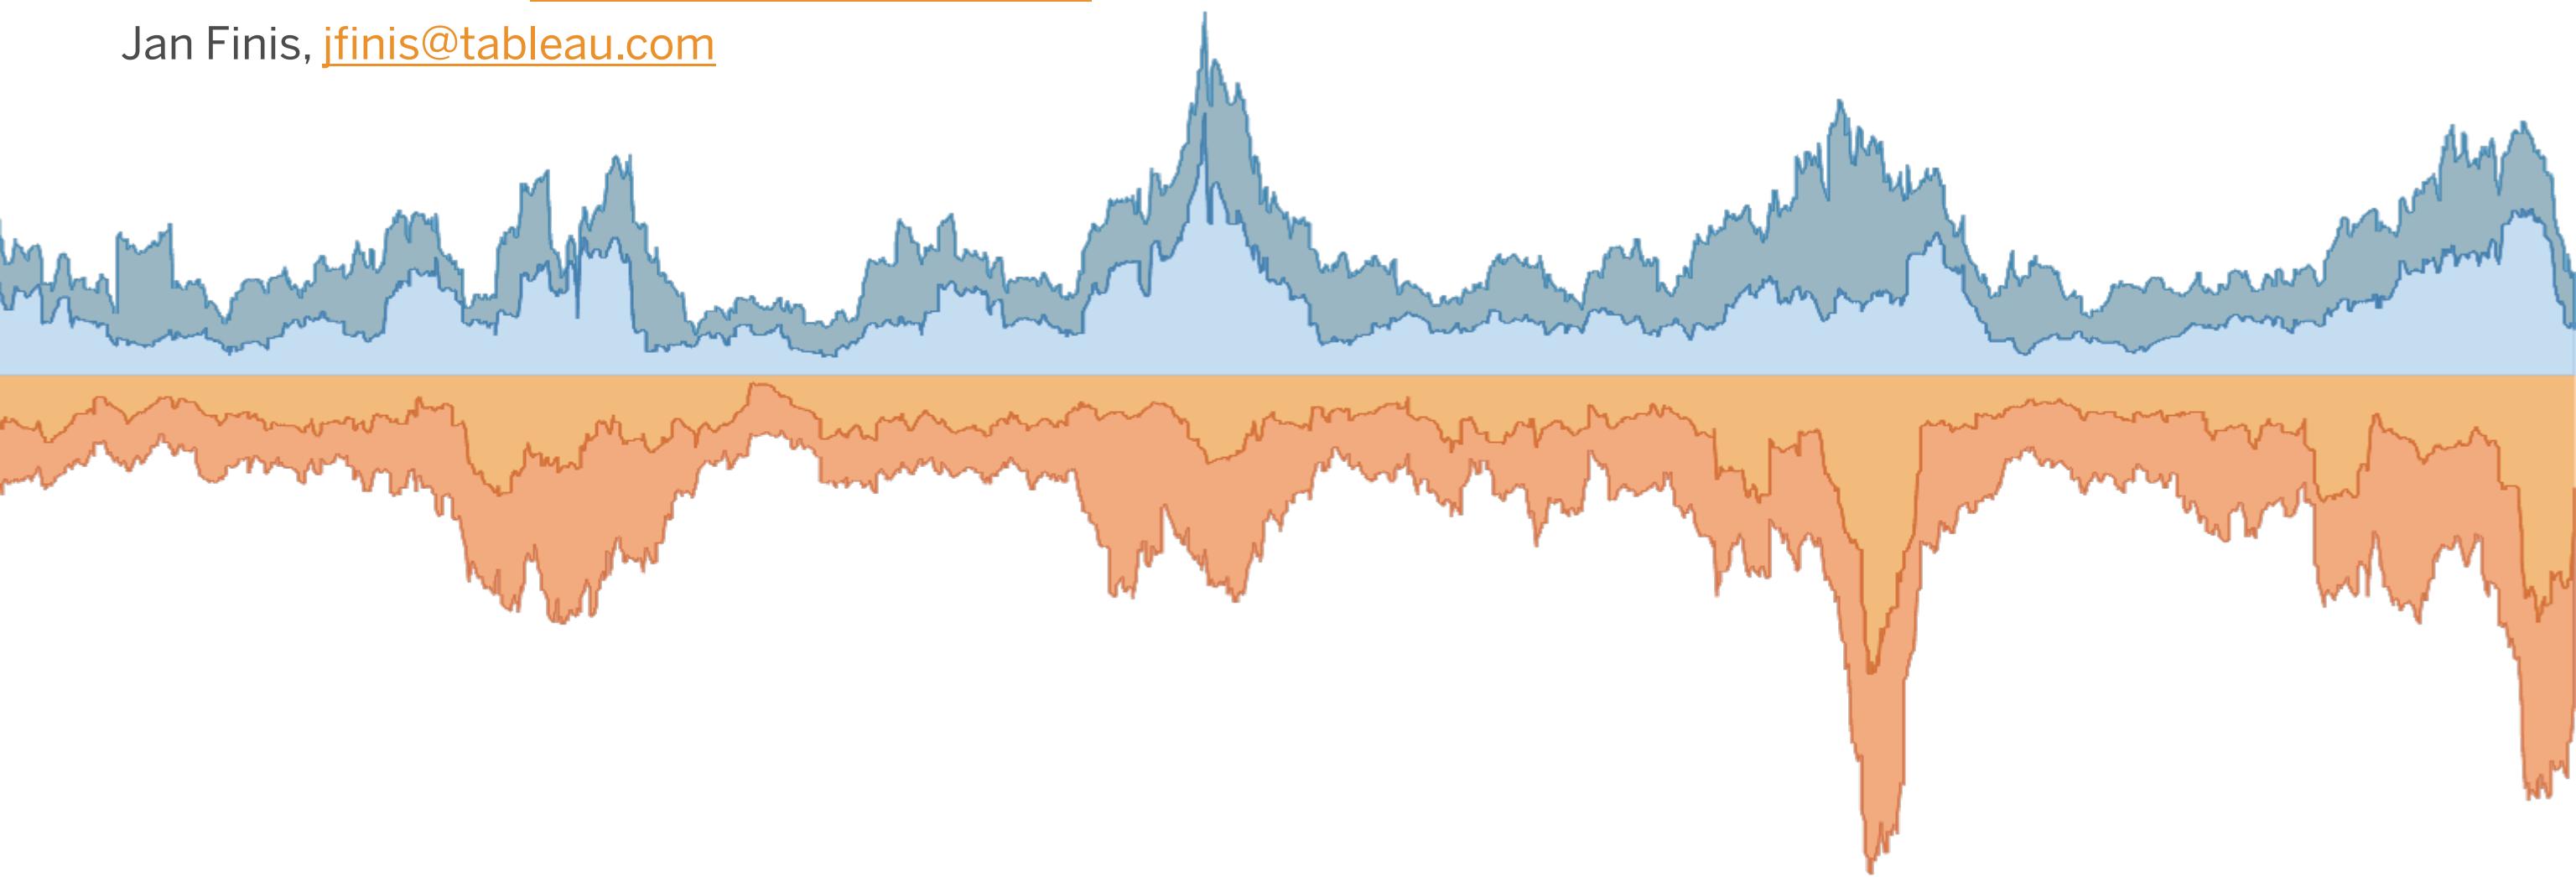


From HyPer to Hyper

Integrating an academic DBMS into a leading analytics and business intelligence platform

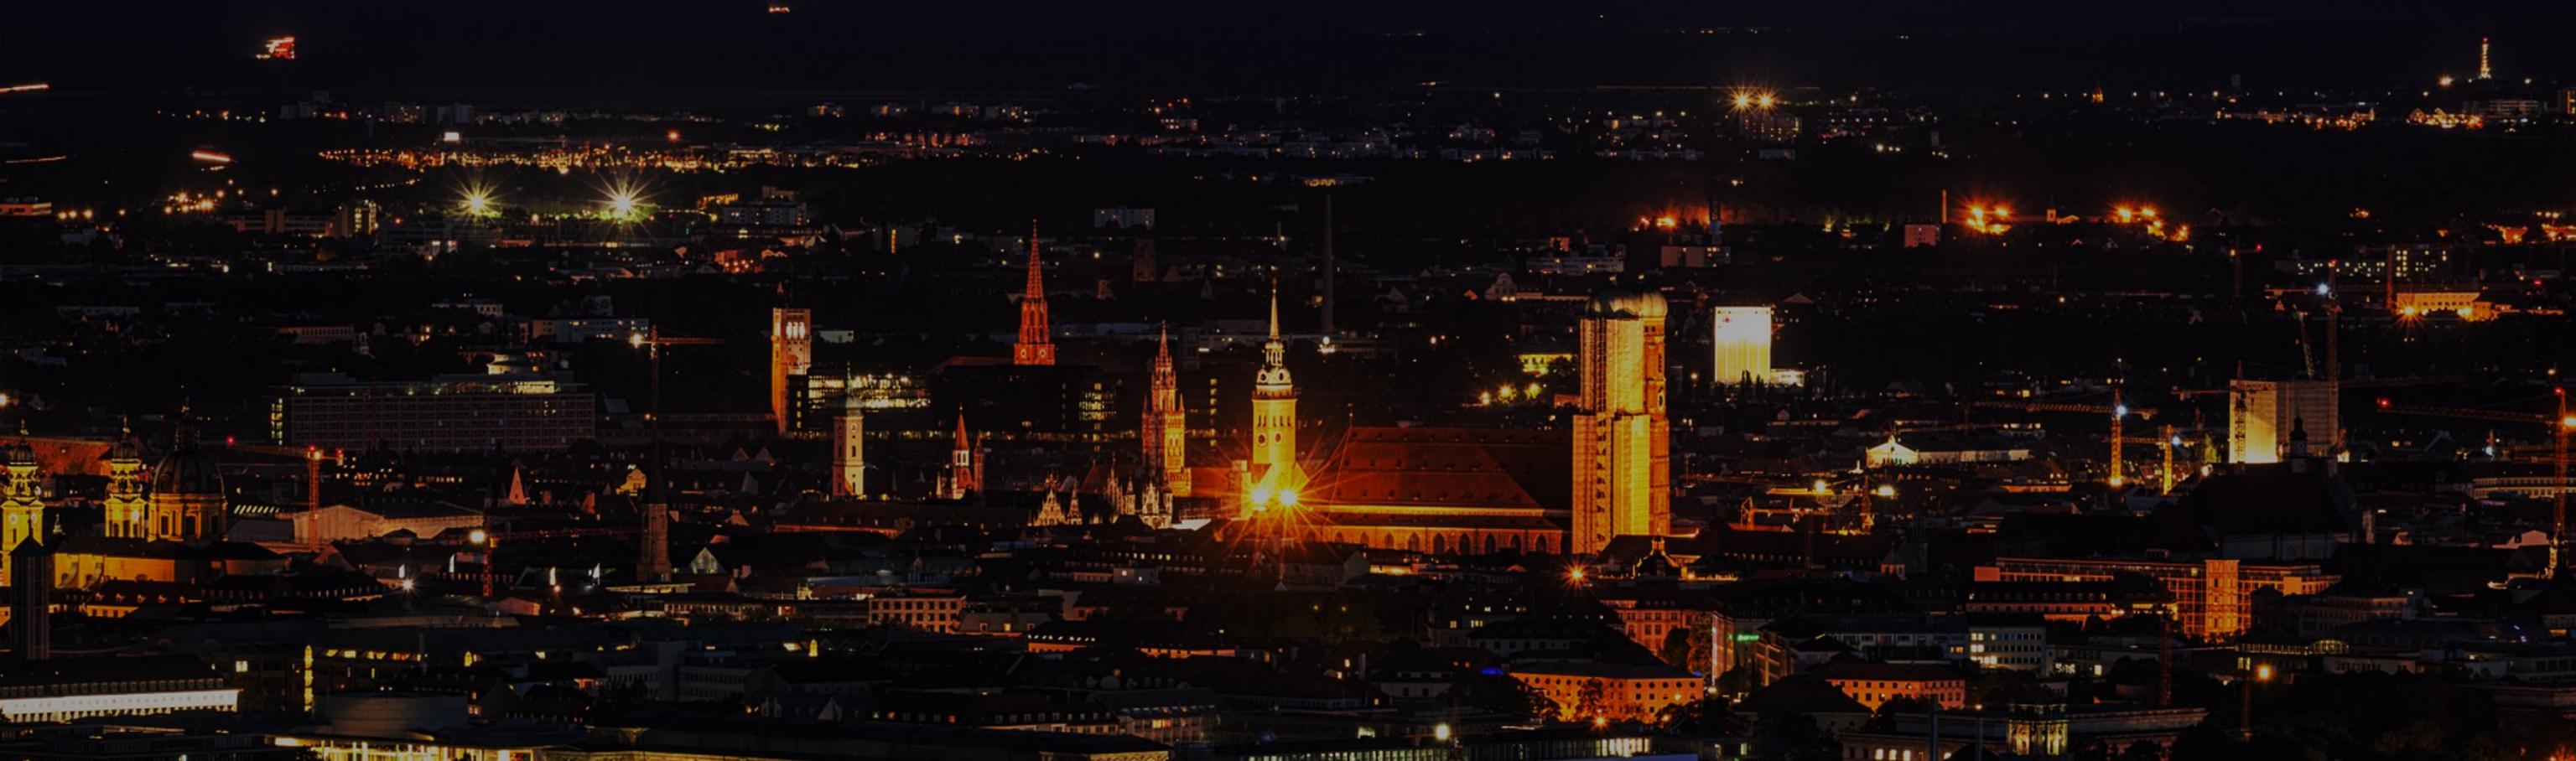
Tobias Muehlbauer, tmuehlbauer@tableau.com

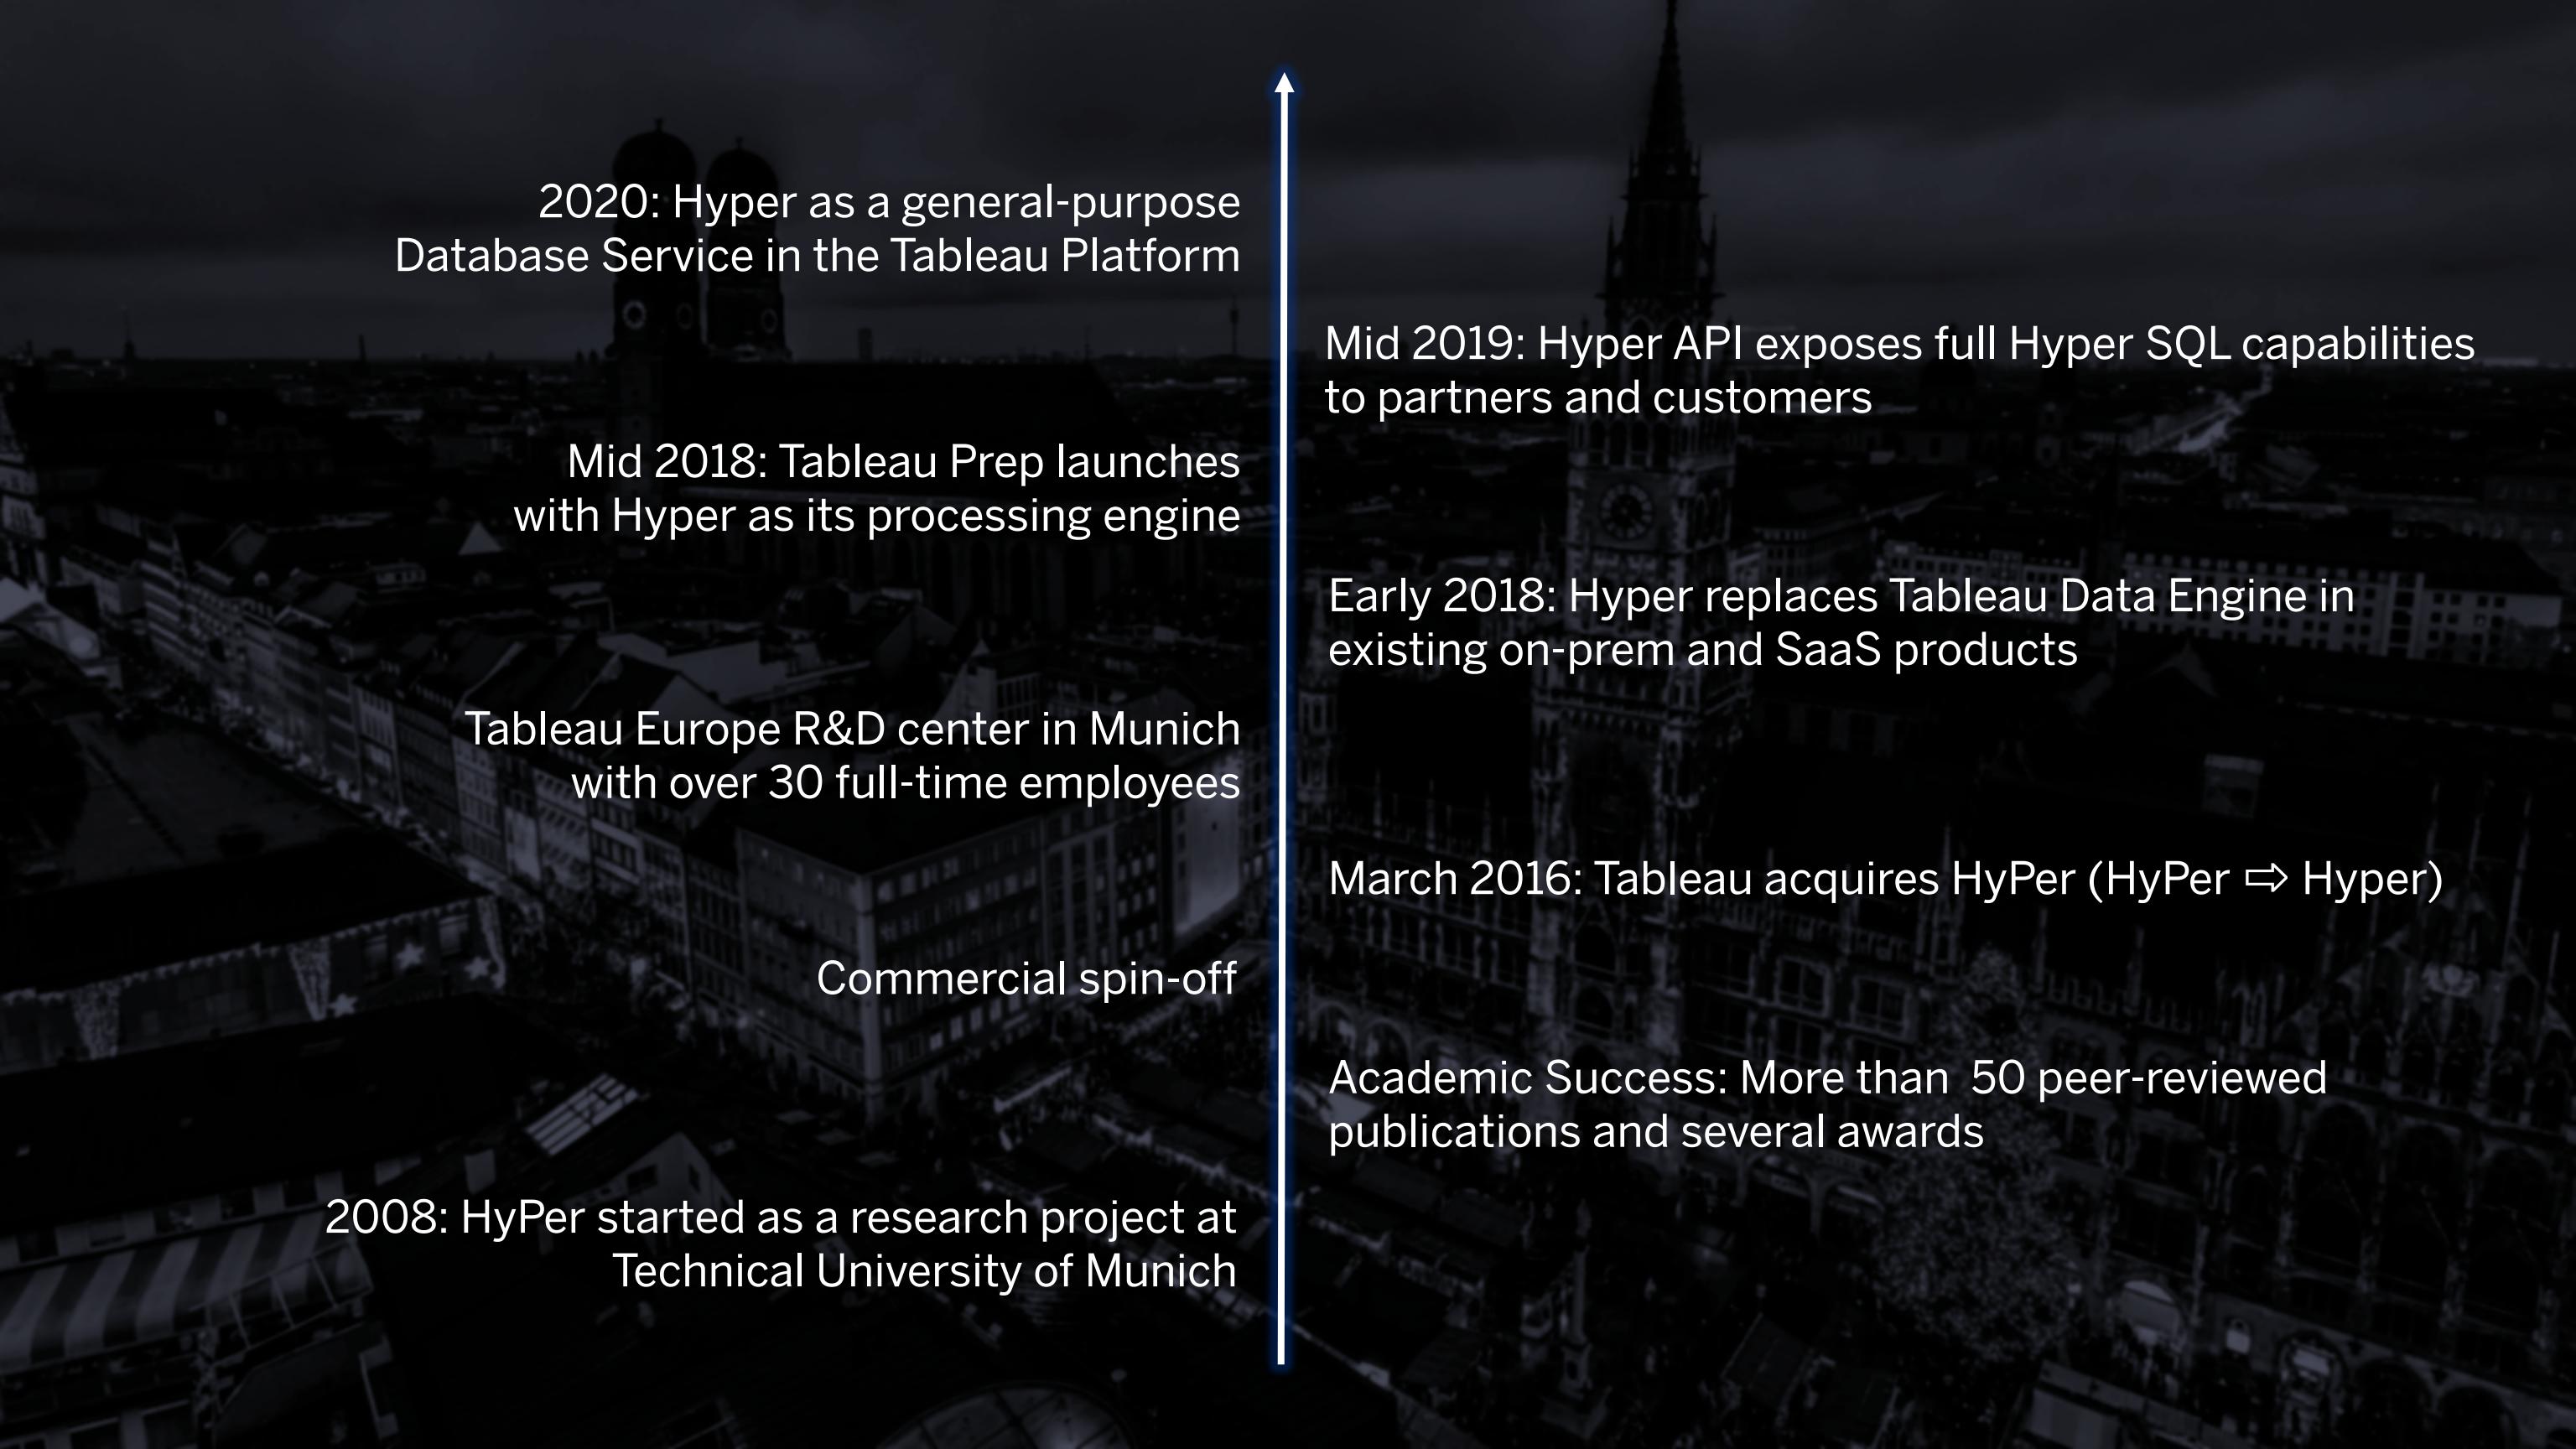
Jan Finis, jfinis@tableau.com





The Story of Hyper





2020: Hyper as a general-purpose Database Service in the Tableau Platform

Mid 2018: Tableau Prep launches with Hyper as its processing engine

Tableau Europe R&D center in Munich with over 30 full-time employees

Commercial spin-off

2008: HyPer started as a research project at Technical University of Munich

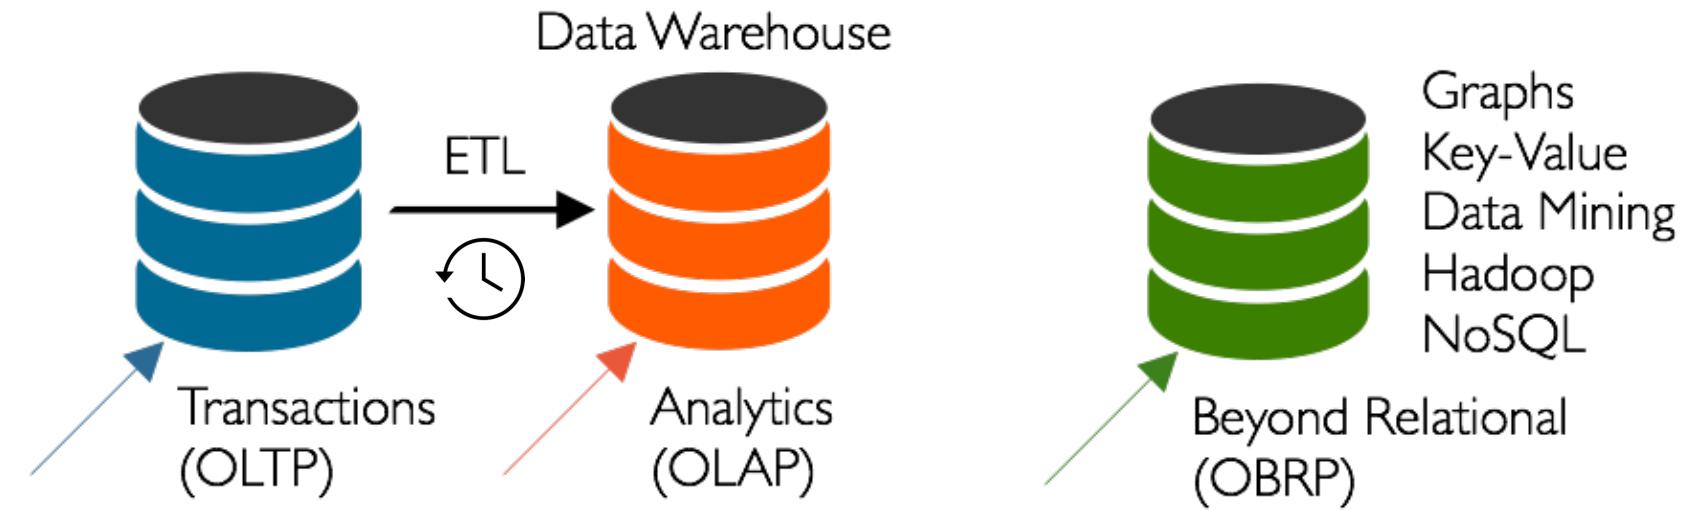
Mid 2019: Hyper API exposes full Hyper SQL capabilities to partners and customers

Early 2018: Hyper replaces Tableau Data Engine in existing on-prem and SaaS products

March 2016: Tableau acquires HyPer (HyPer \Rightarrow Hyper)

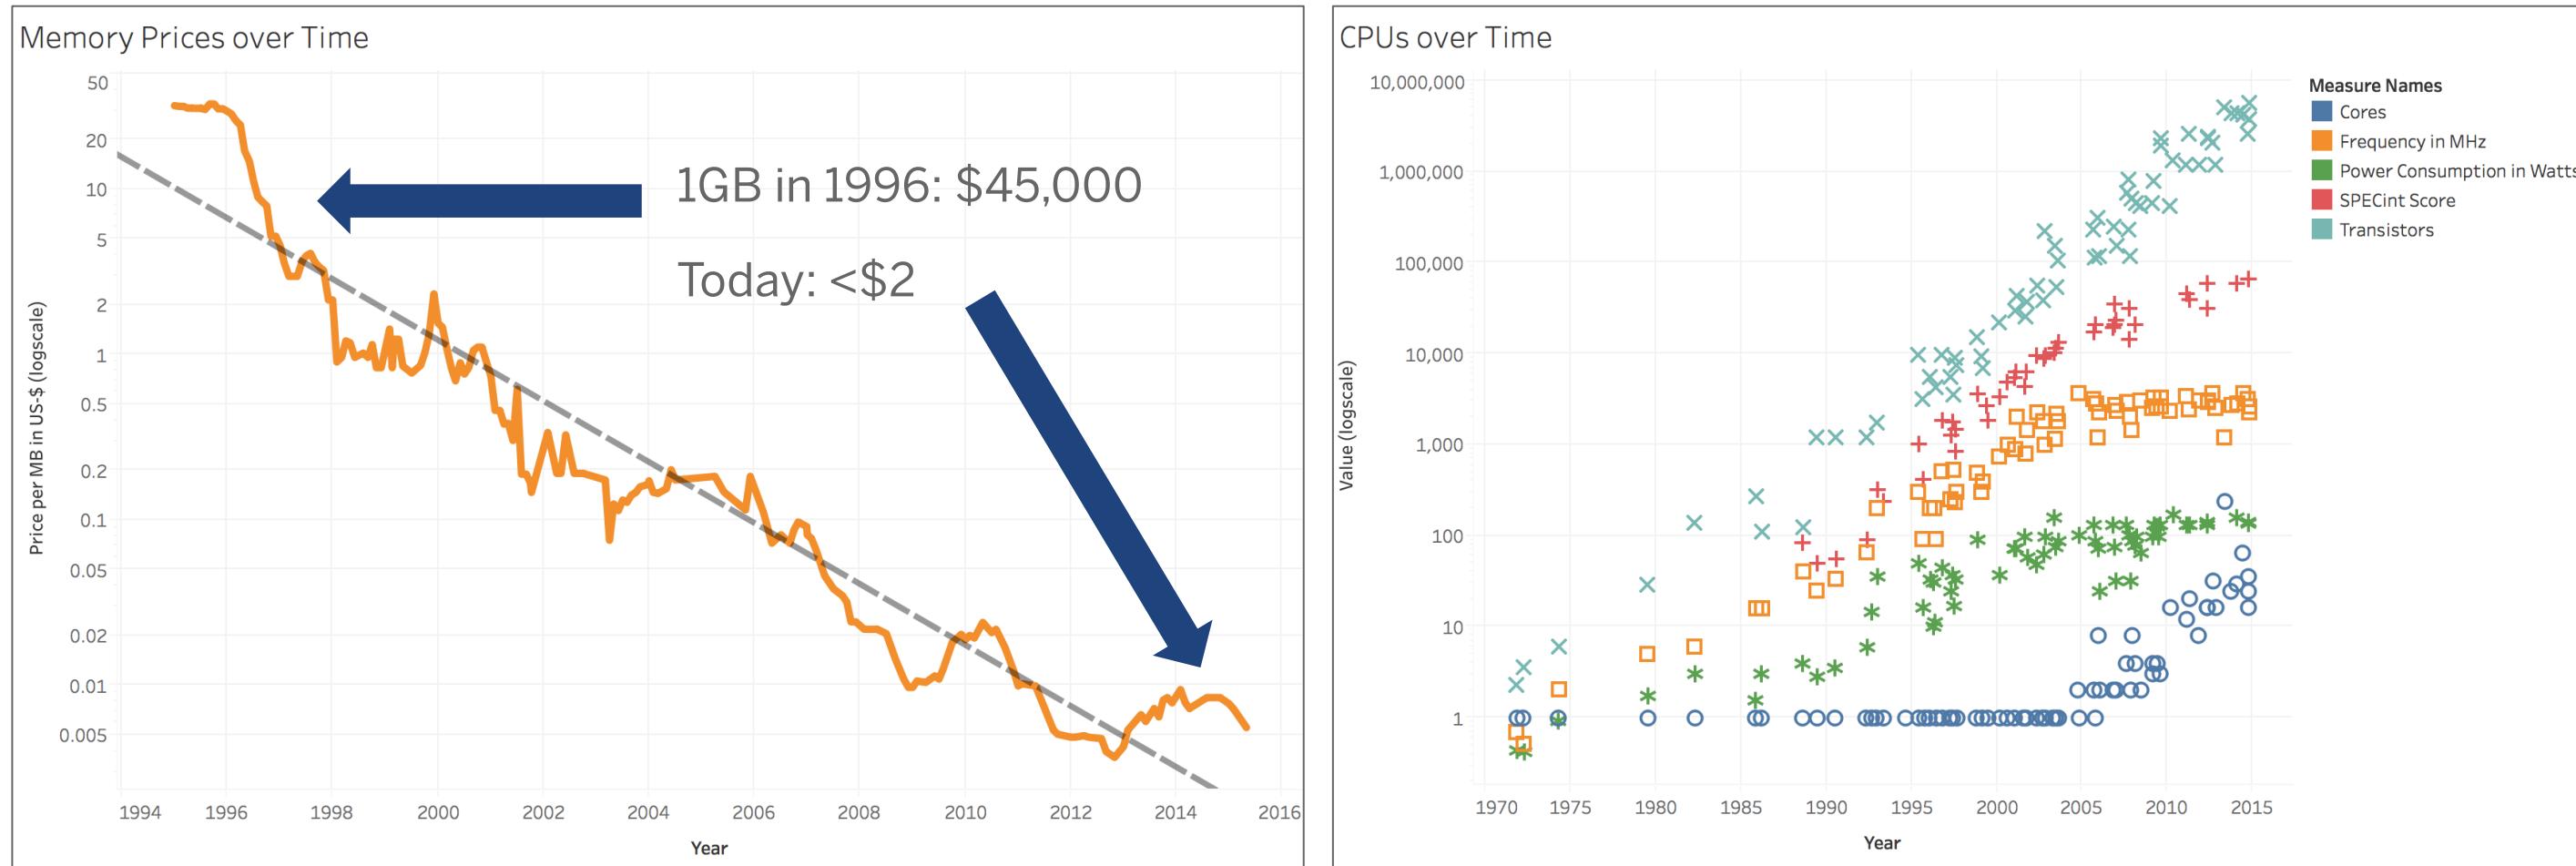
Academic Success: More than 50 peer-reviewed publications and several awards

Context: One Size Fits All?



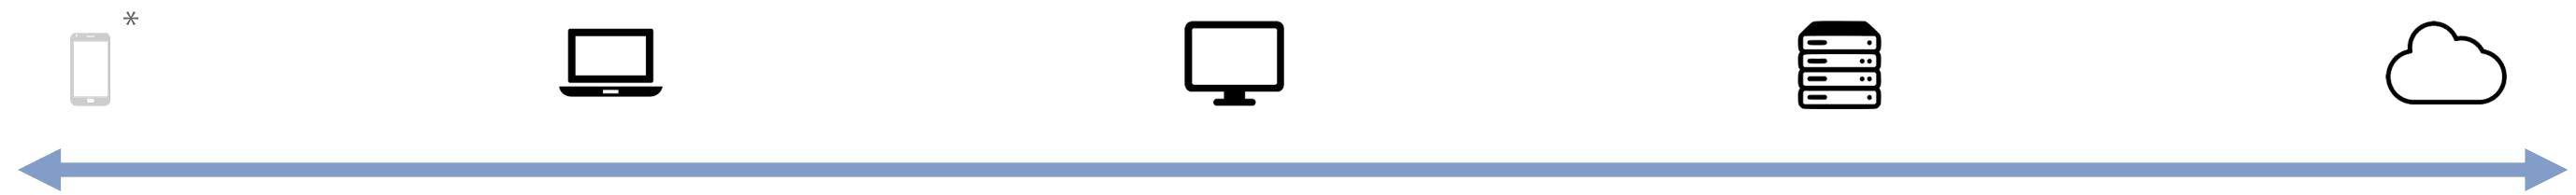
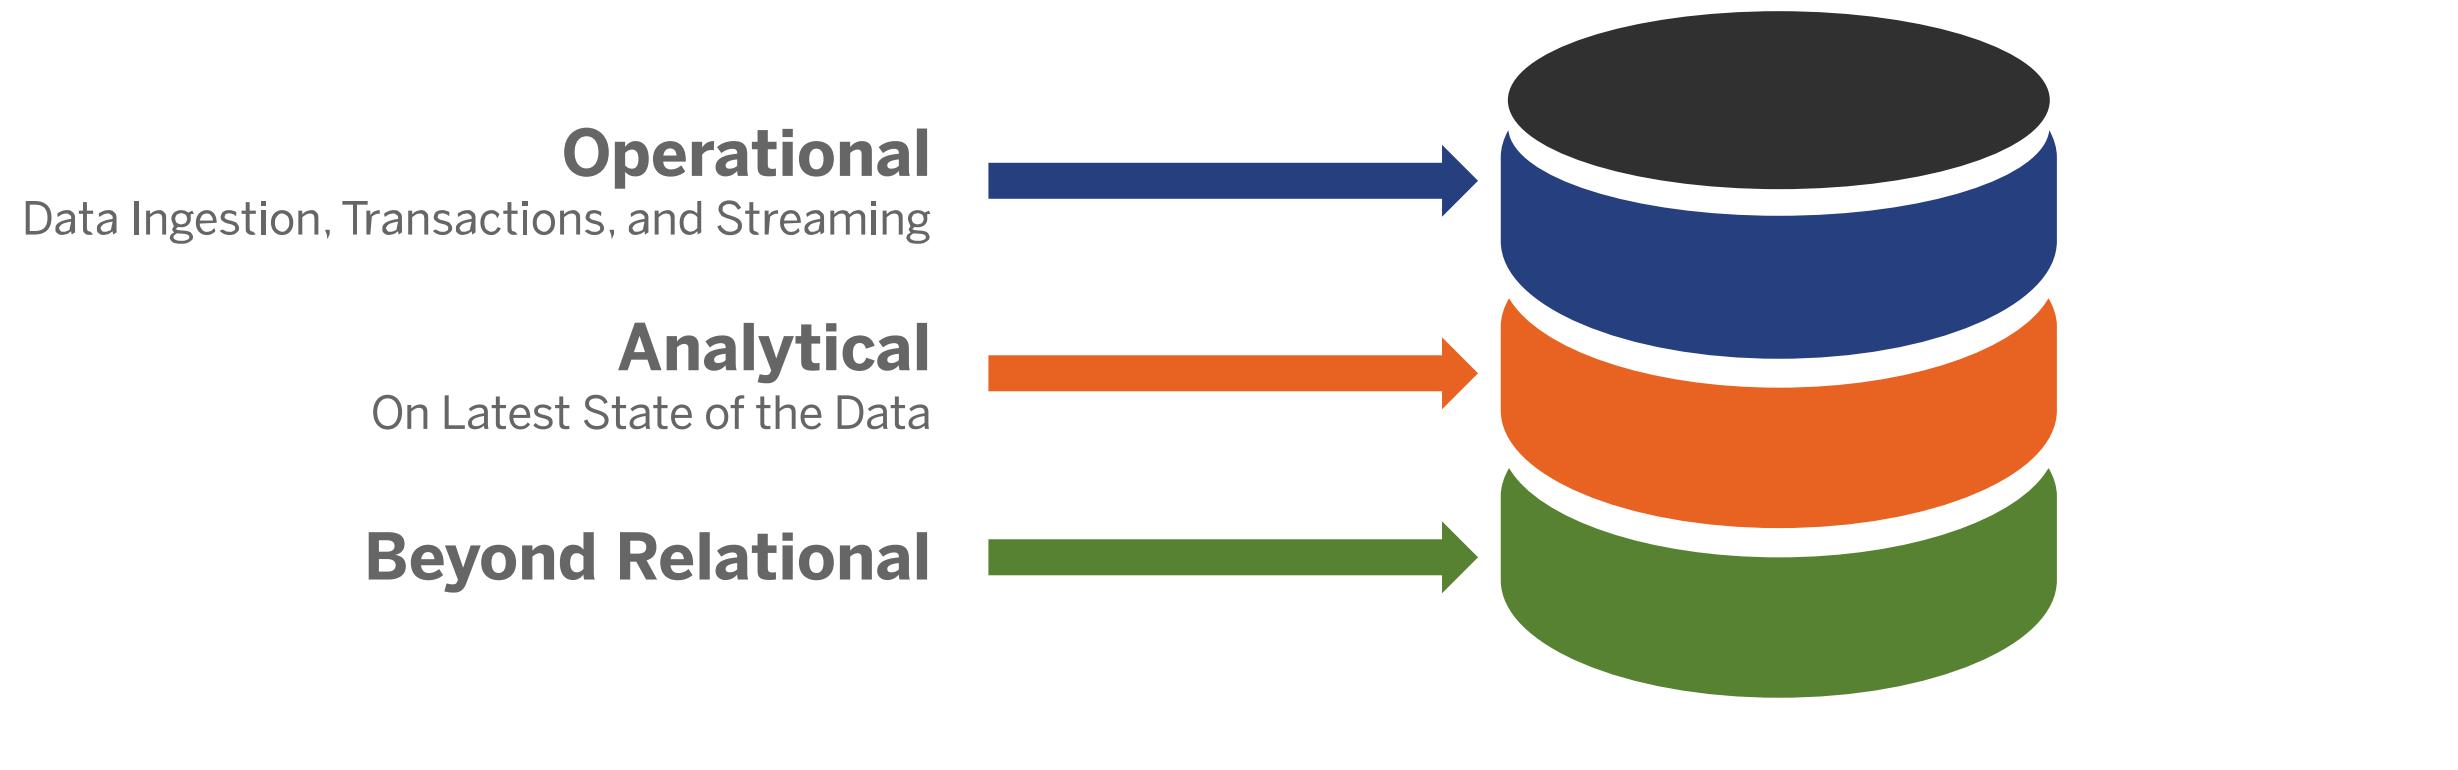
One size does **not** fit all, but what about data lag and disconnect?

Context: A Changing Hardware Landscape



In order to leverage modern hardware, **databases need to change**.

The Idea Behind Hyper



*Prototype

Optimized for Underlying Hardware & Operating System





One System

Hyper is developed as a **general-purpose database system** that combines **transaction processing, data ingestion, and data analytics**.



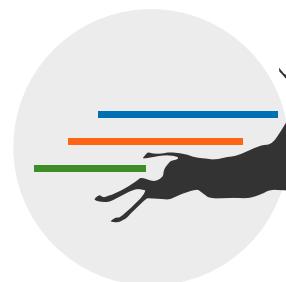
One State

Transaction processing, data ingestion, and data warehousing all on the same state to **enable real-time analytics of the latest state of your data**.



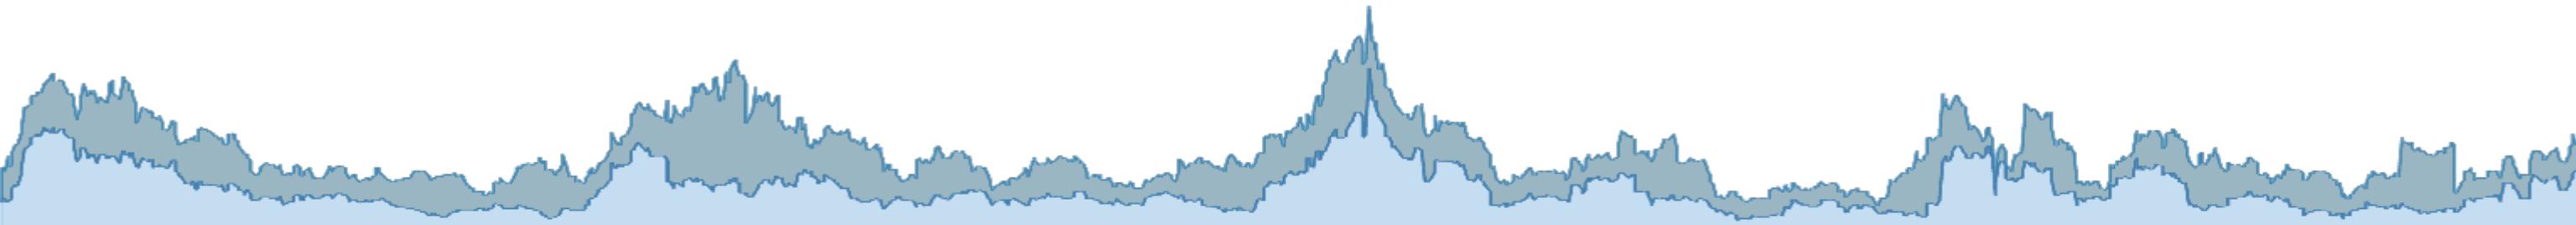
No Tradeoffs

Hyper makes no tradeoffs when it comes to **ACID guarantees** and **SQL-92+ language support (based on PostgreSQL dialect)** that stood the test of time.

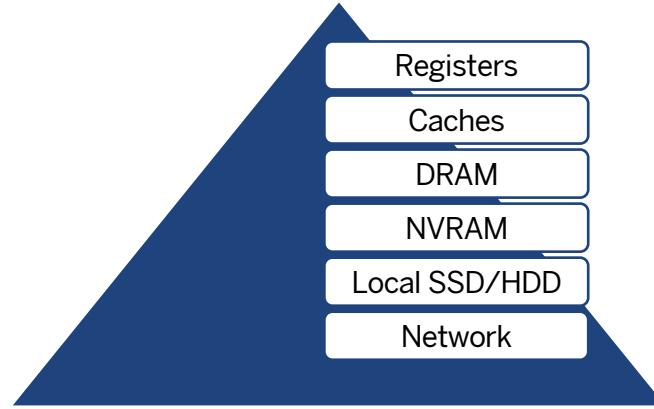


No Delays

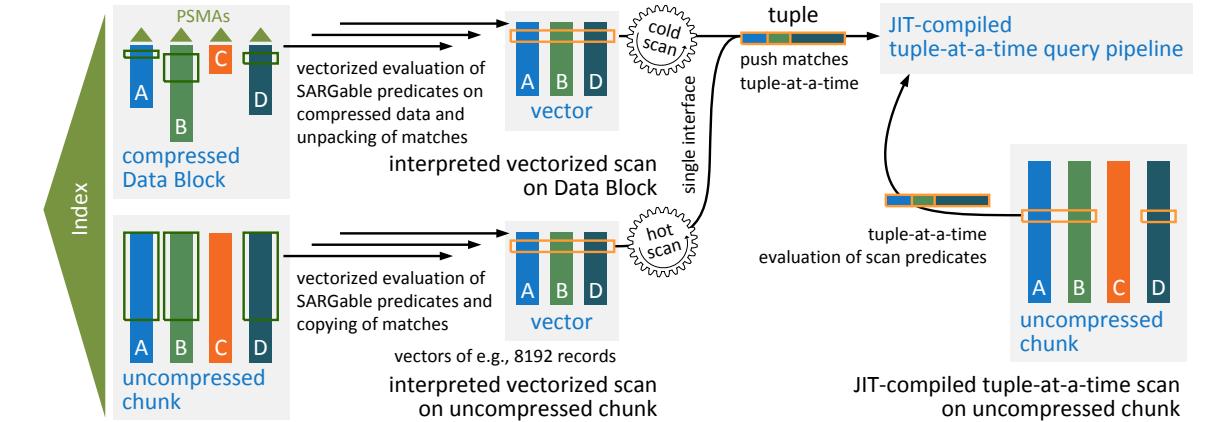
Hyper **scales with available hardware resources** to allow **highest performance on all workload classes**.



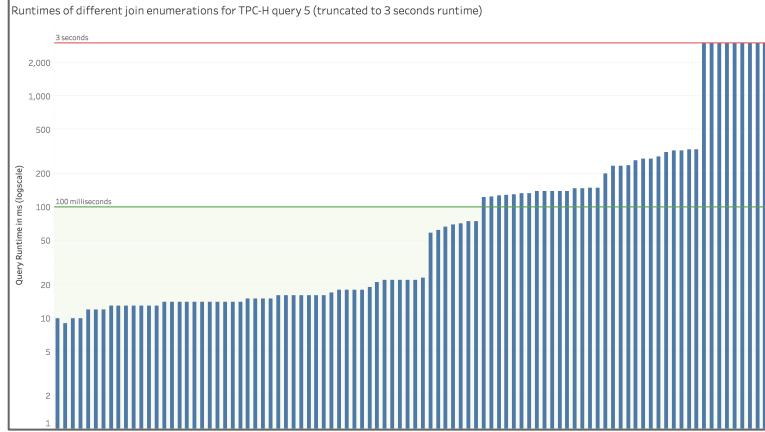
Inside Hyper



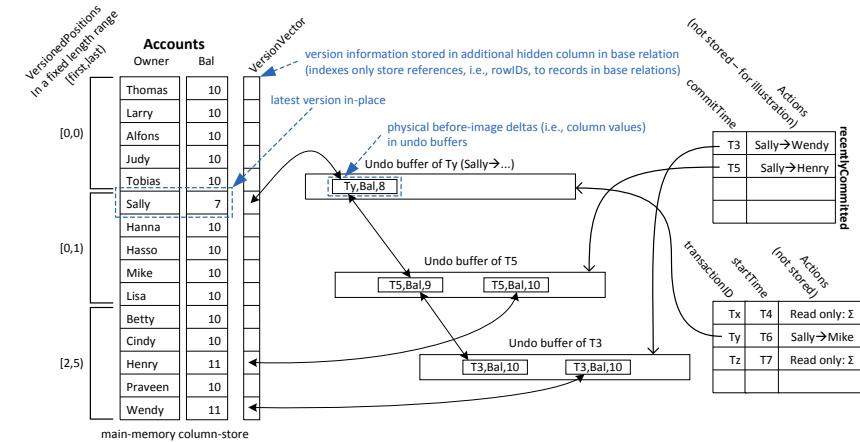
Optimized for Storage Hierarchy



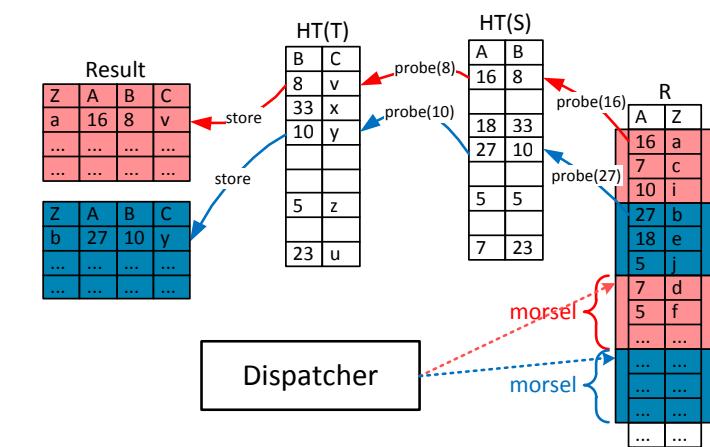
Query Compilation & Vectorized Scans



Advanced Query Optimization



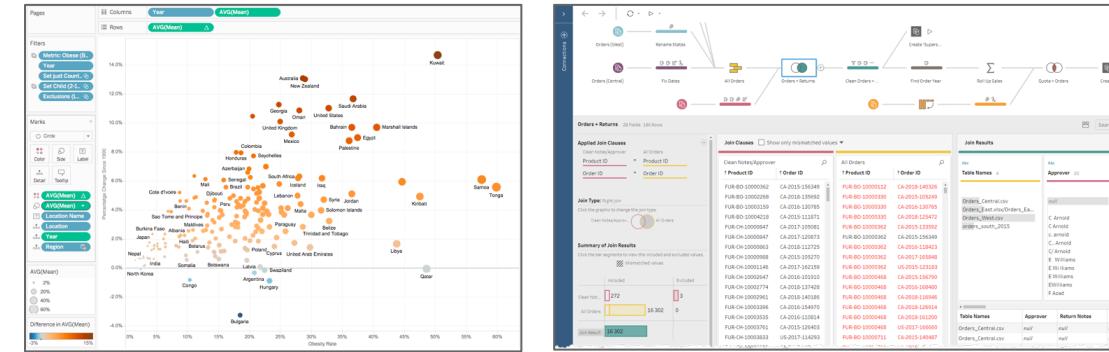
Fast MVCC



Morsel-Driven Parallelization

Hyper in Tableau

Desktop Online Public Server Prep



SQL Queries



Results



Hyper

(Bulk) Insert, Update, Delete, Streams, Files

Extract

Hyper API

Federation

Prep

From HyPer to Hyper: Challenges

Support

- Limited support provided up to 30 months after major product version release
- Compare performance across releases and database engines
- Semantic differences

Infrastructure

- Windows, Linux, and macOS
- Small laptops to large-scale servers and Cloud deployments

Workload

- Long tail of query complexity generated by Tableau
- Wide variety of data set characteristics

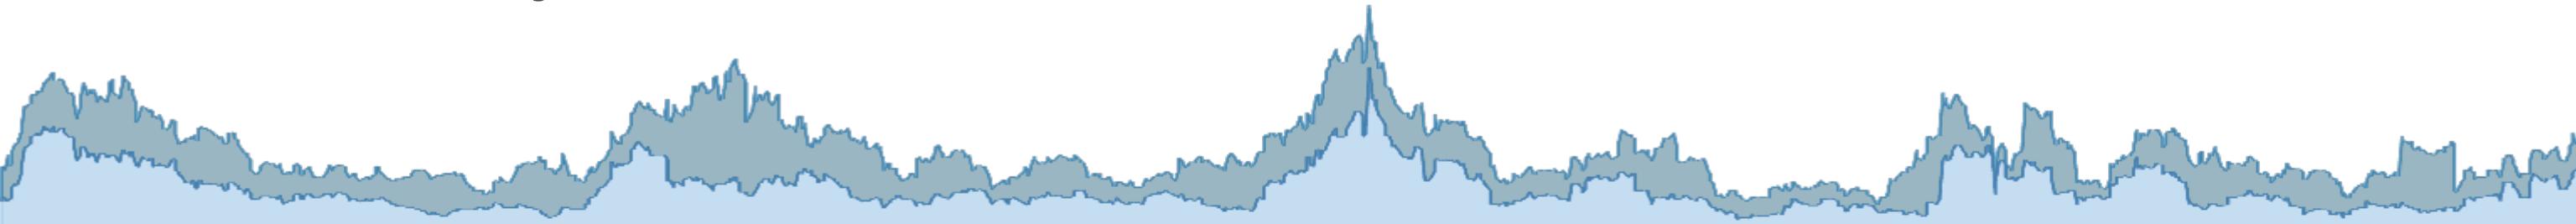
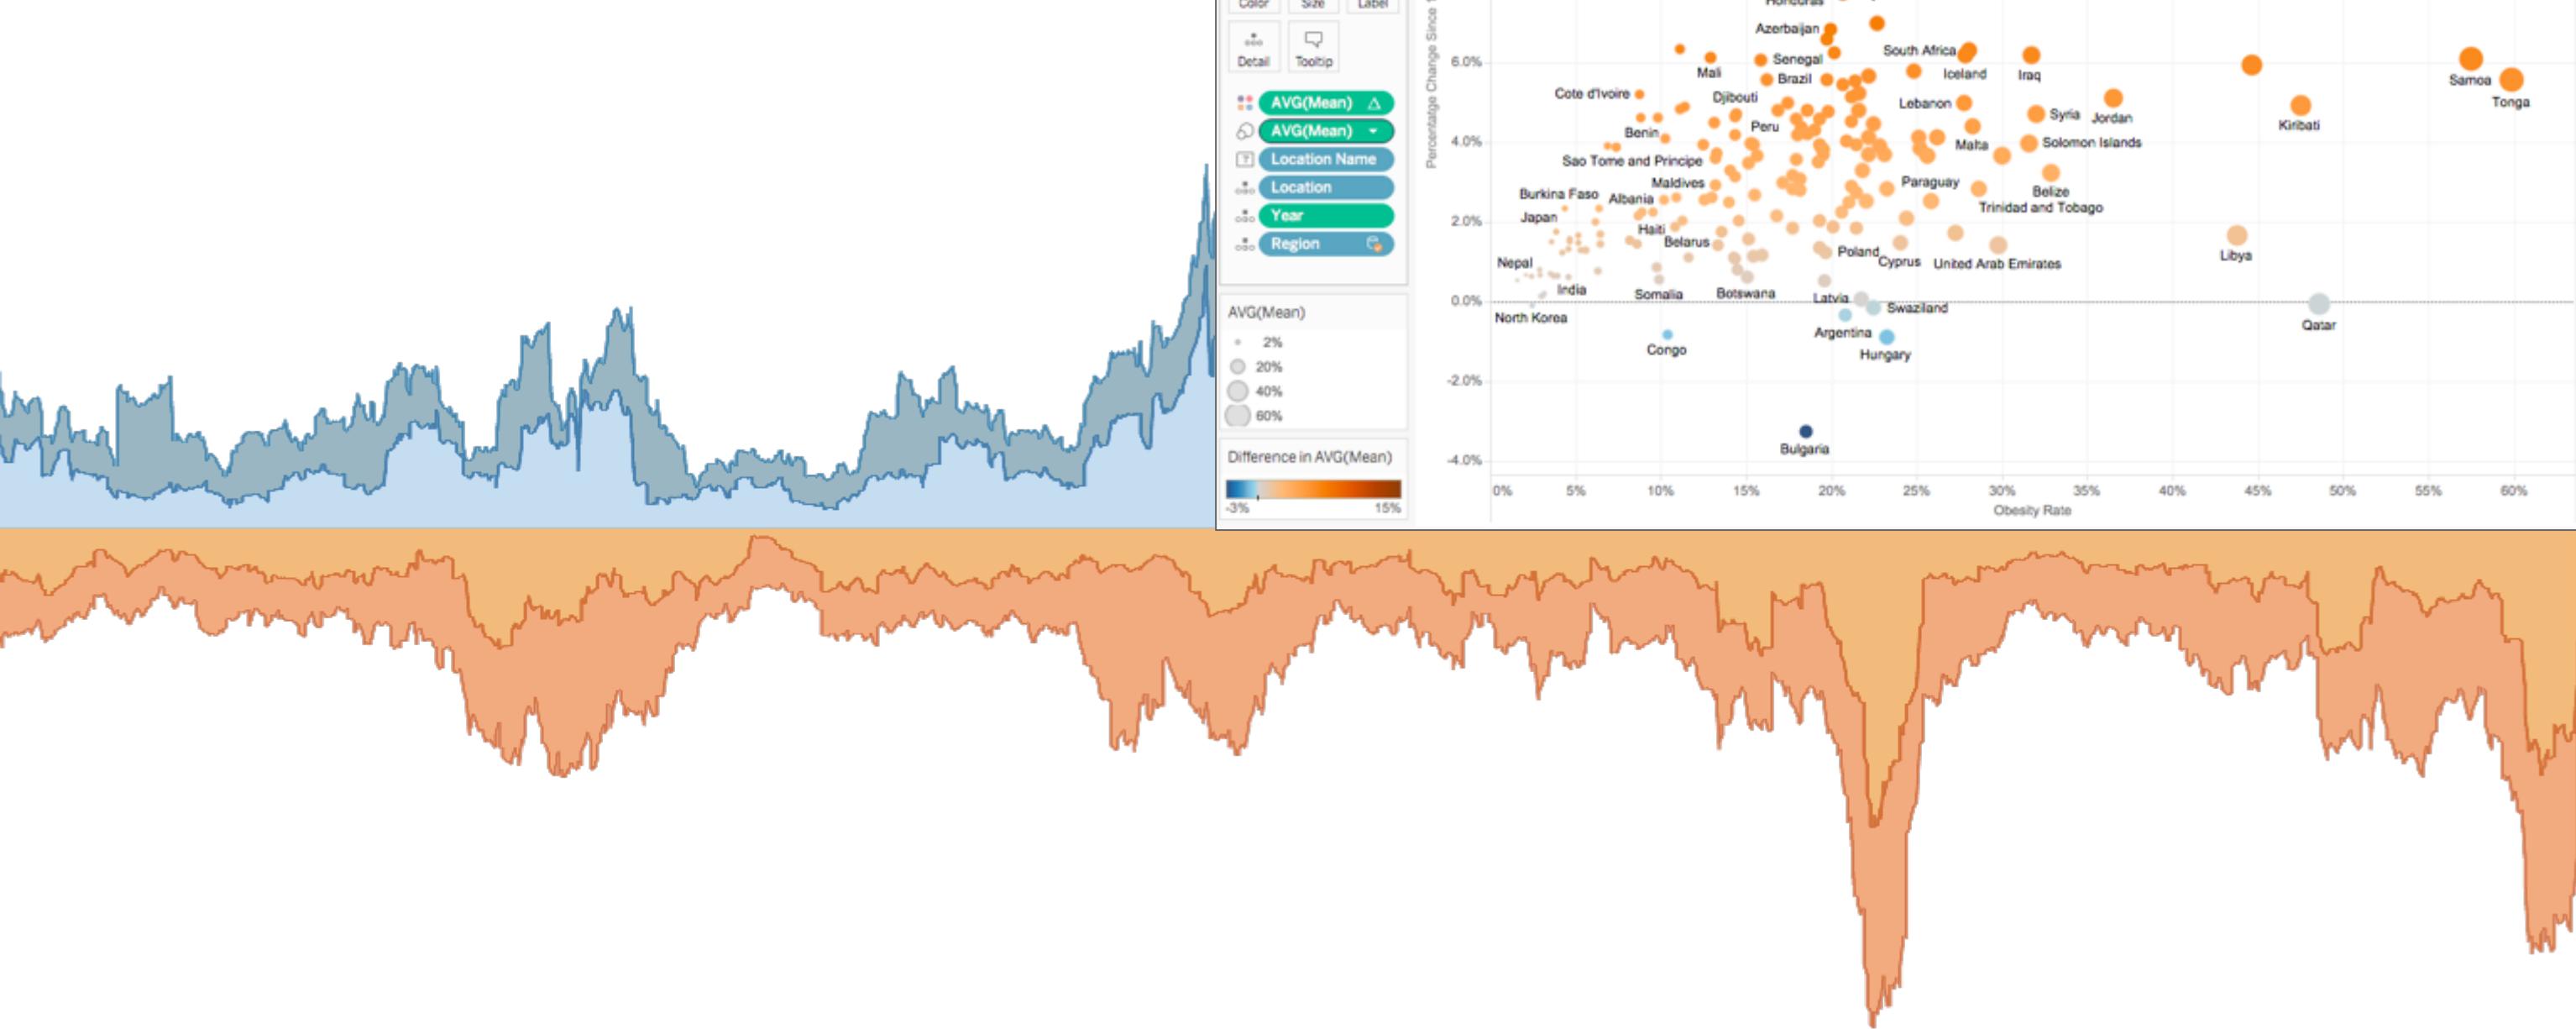


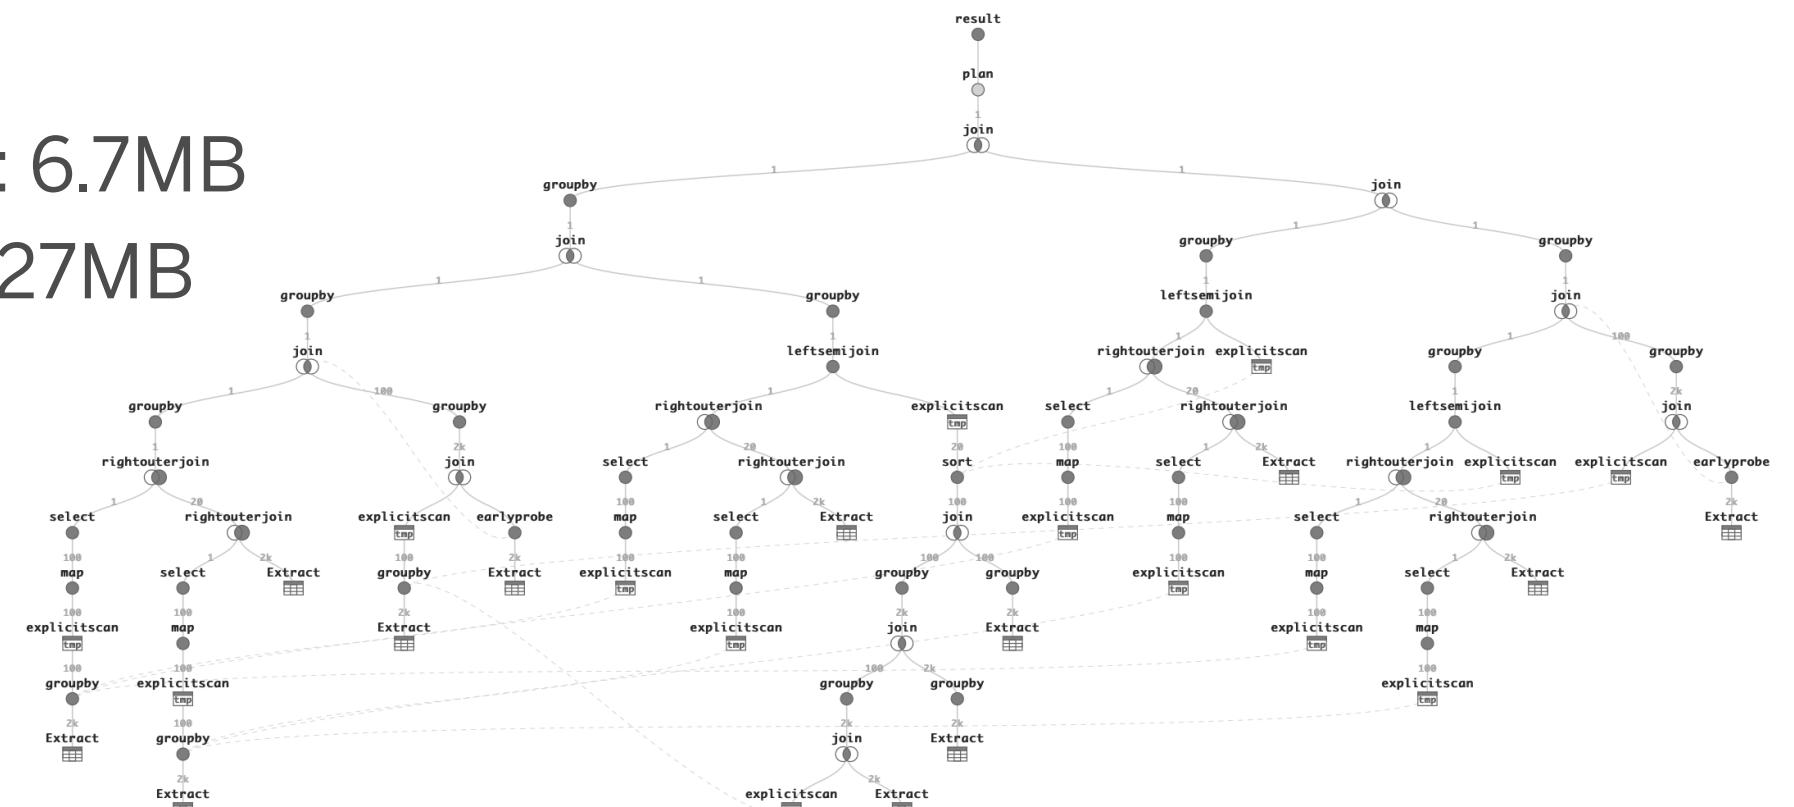
Tableau Workloads

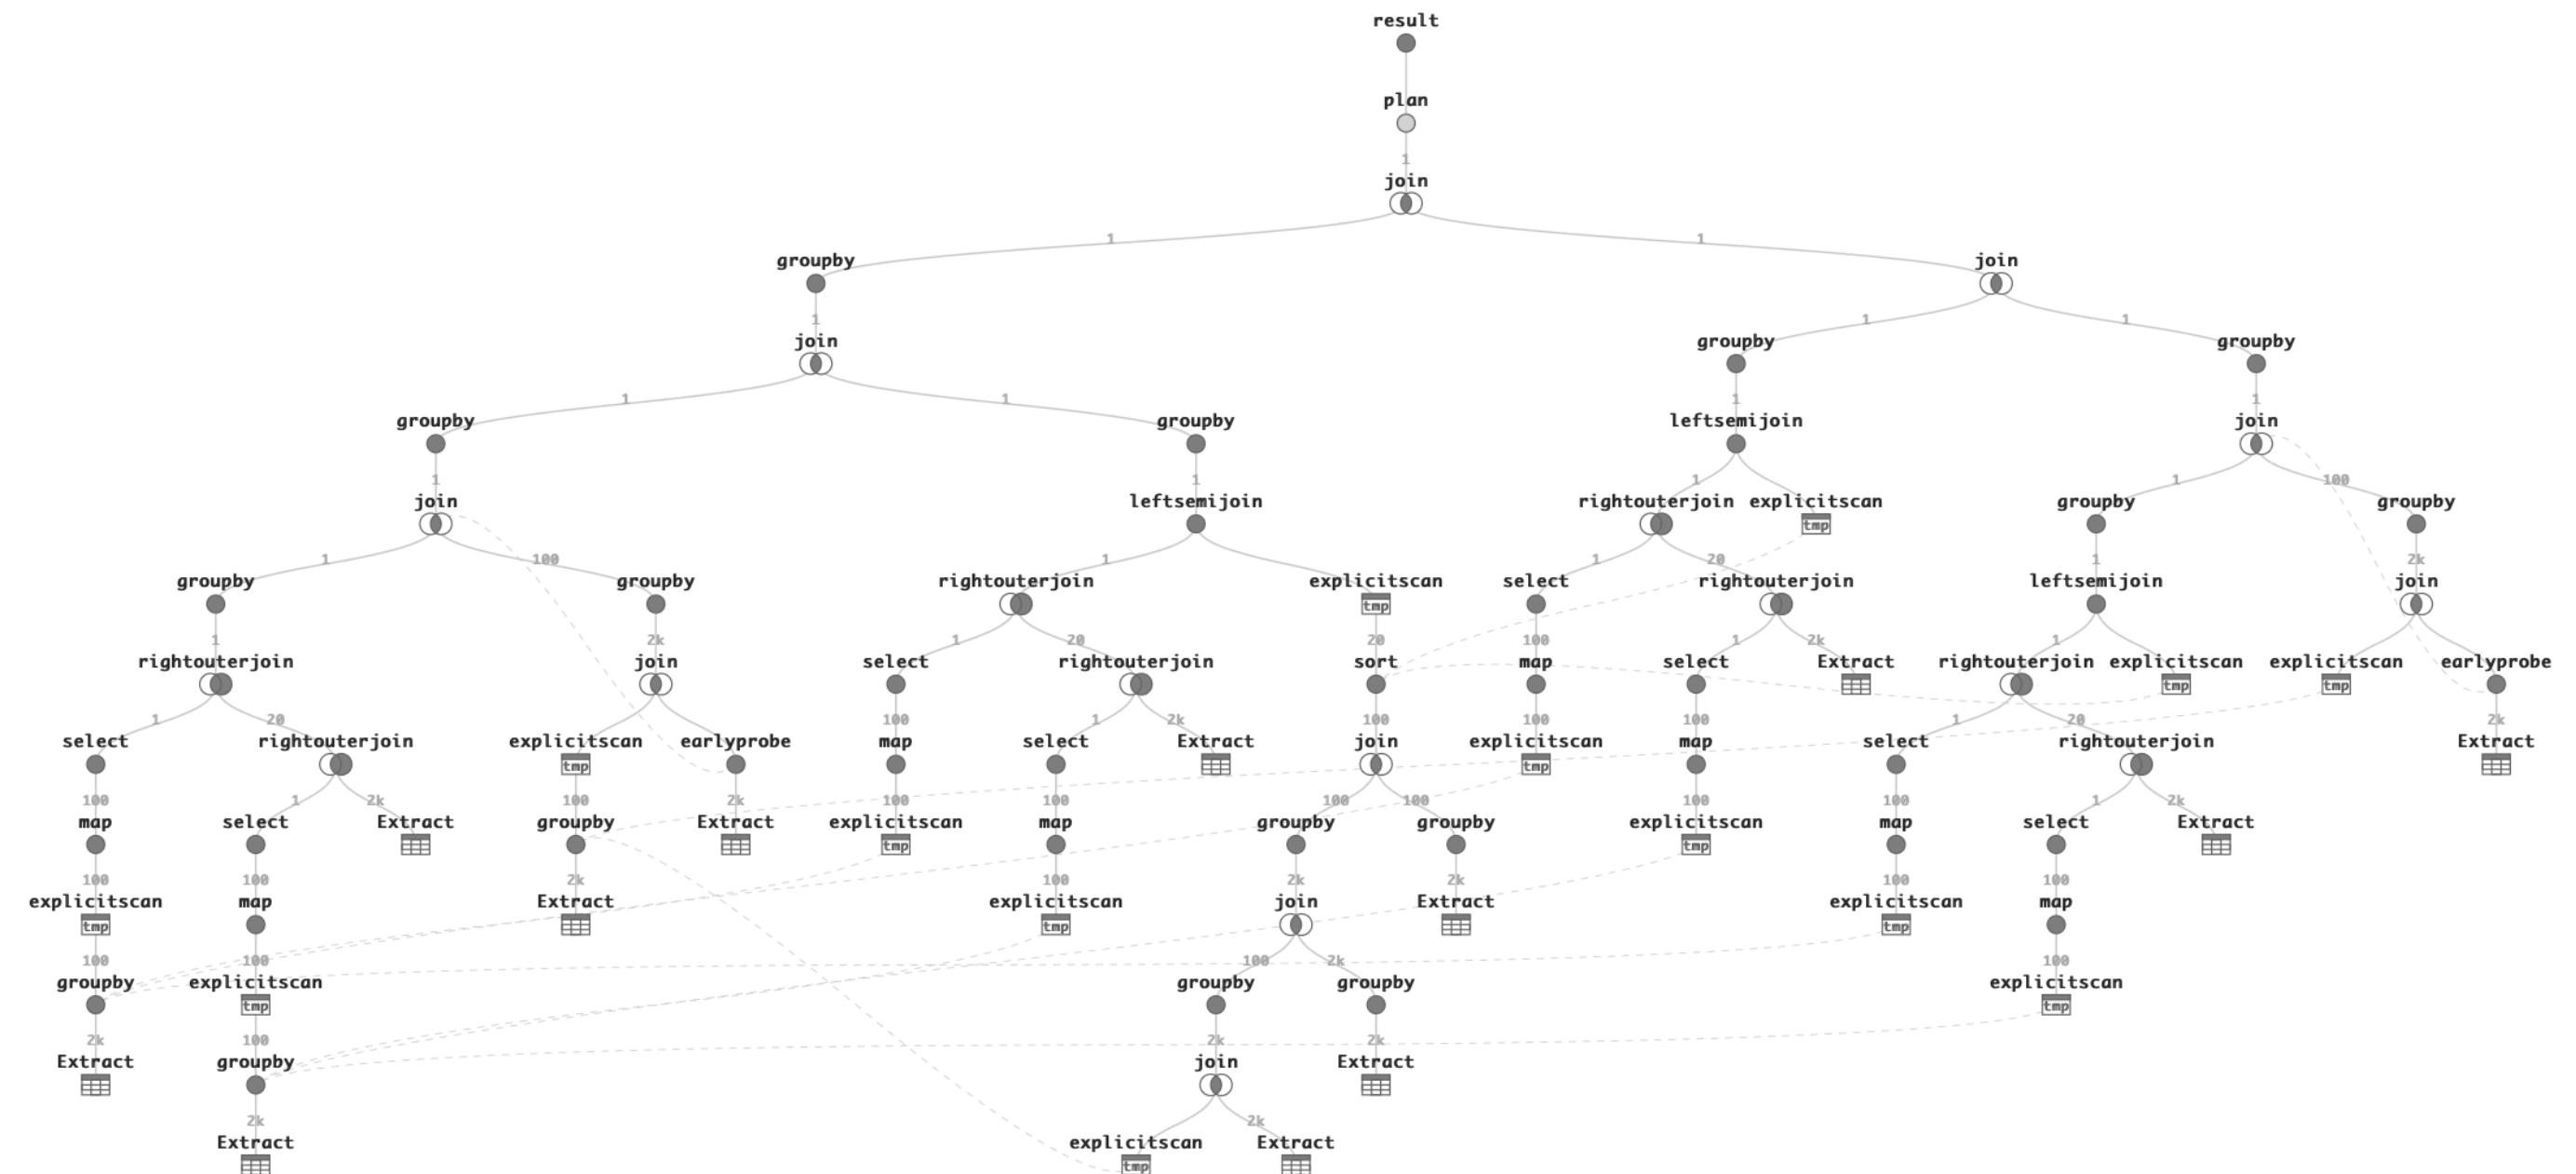


What Tableau Workloads Look Like

- Most queries are “small”: Only 0.5% larger than 5KB SQL Text
- But: **Huge** outliers
- Largest query in our data set: 6.7MB
- Largest query we saw so far: 27MB

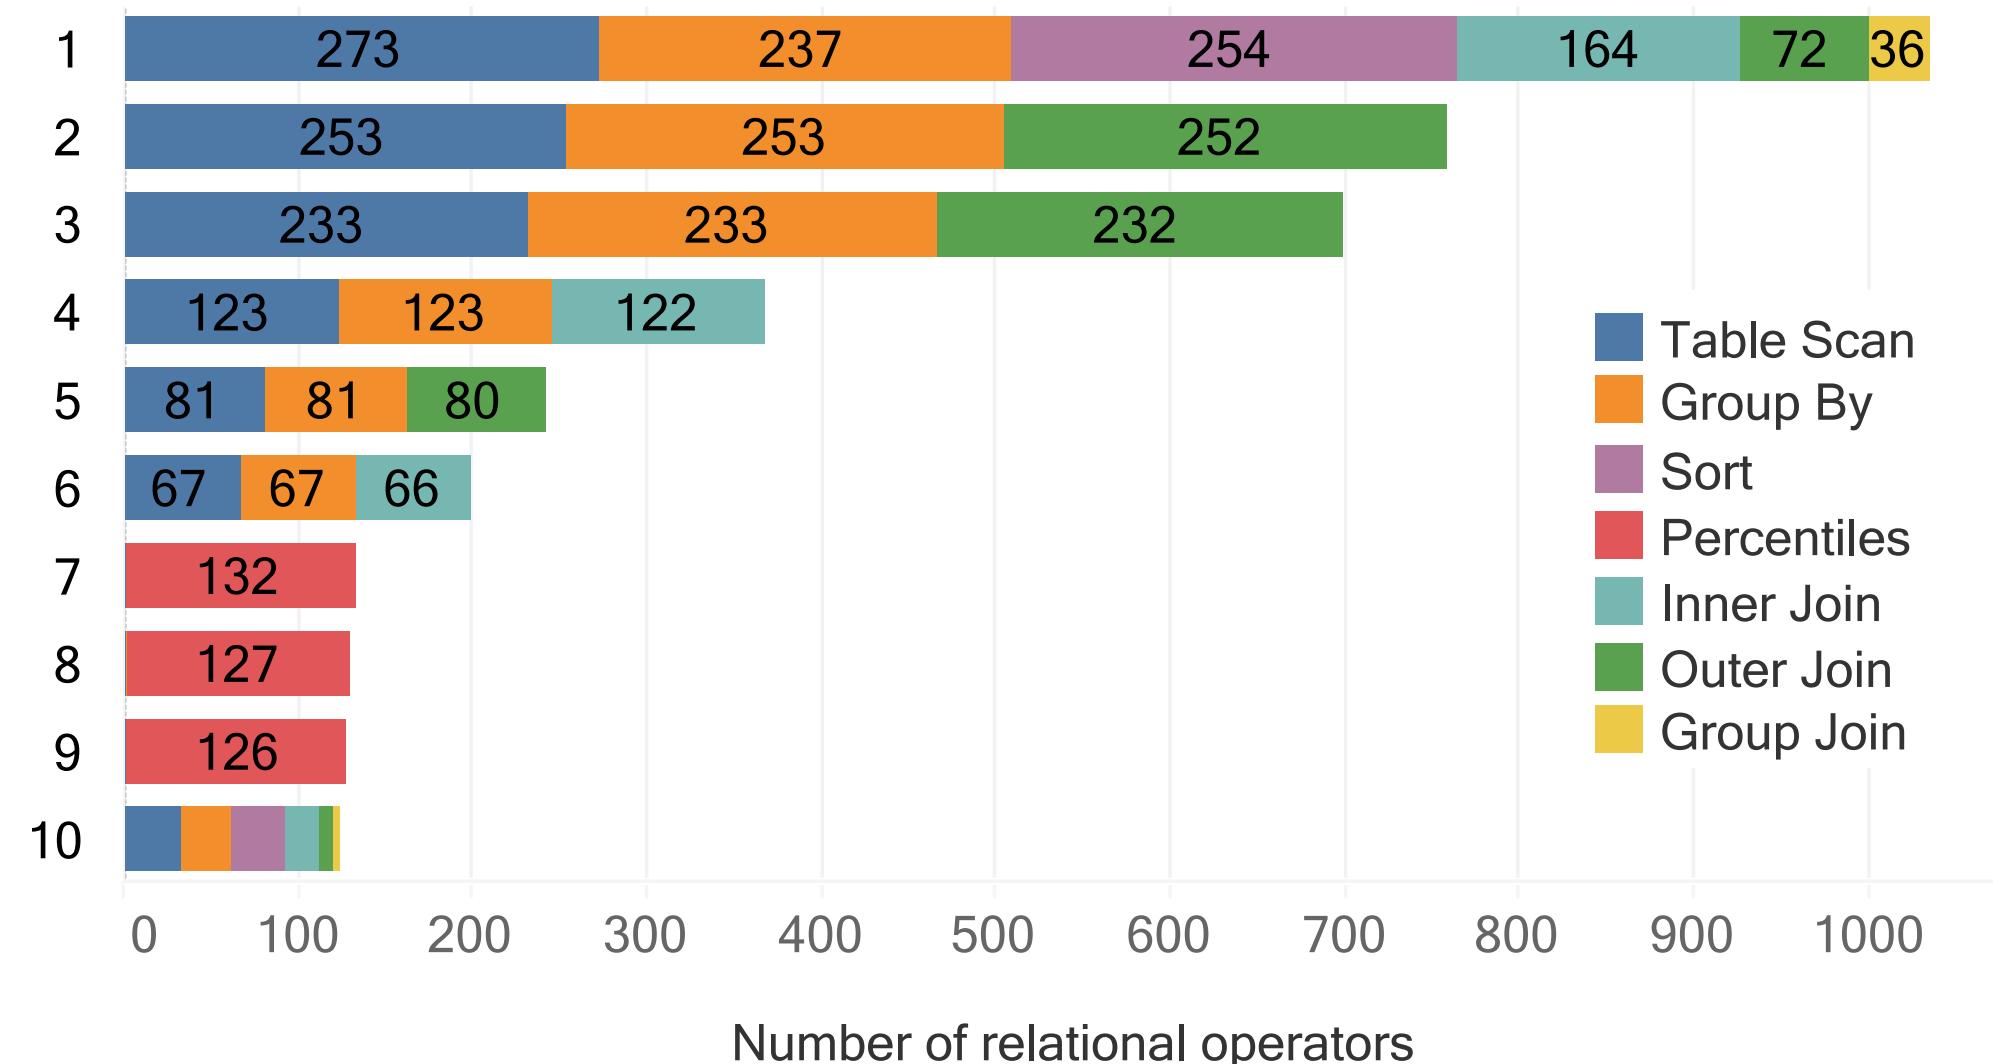
And that's not all due
to constant strings...





Need a query plan visualizer? <https://github.com/tableau/query-graphs/>

What Tableau Workloads Look Like



Vogelsgesang et al.: Get Real: How Benchmarks Fail to Represent the Real World. DBTest'18

Replacing Tableau's Old Data Engine



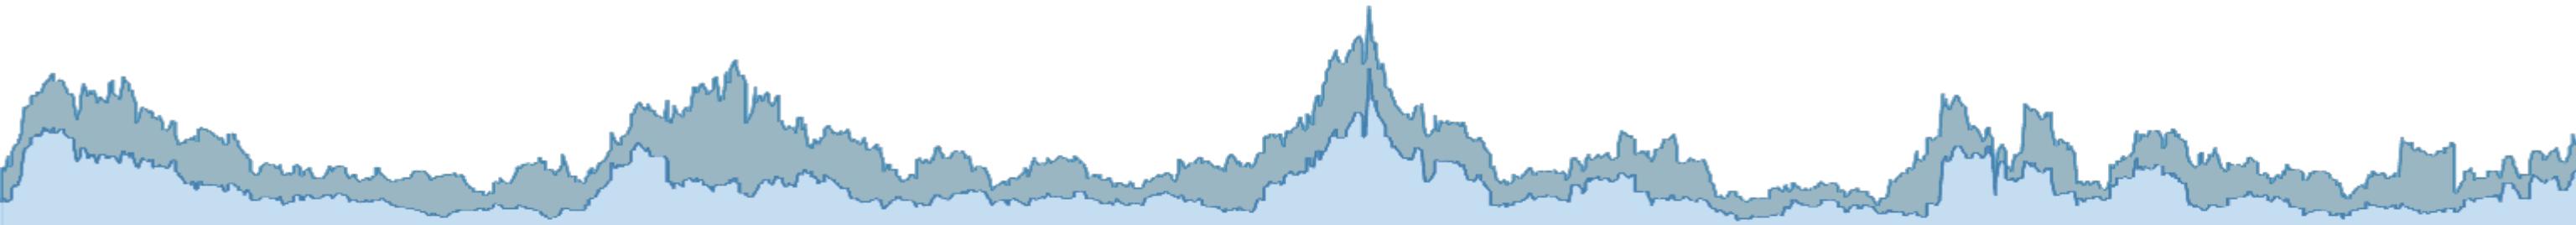
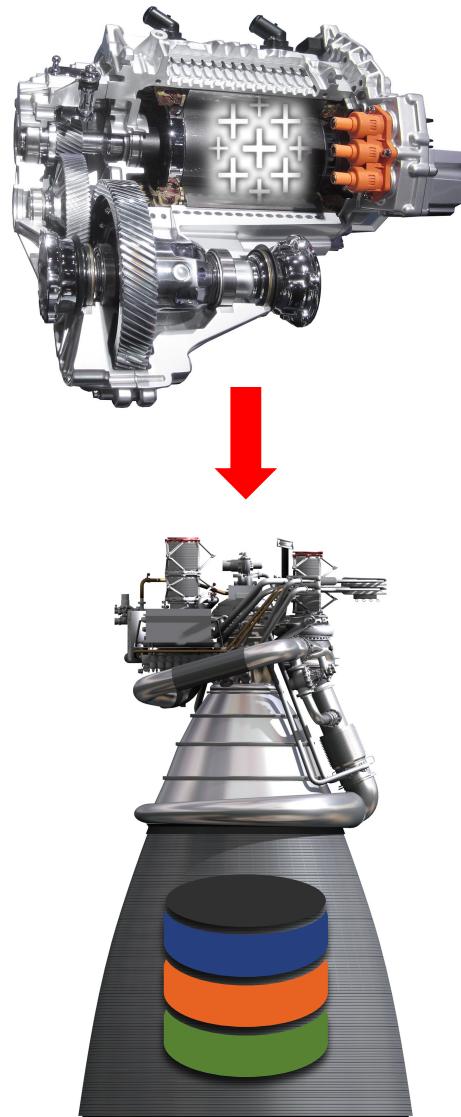
Replacing Tableau's Old Data Engine

Tableau's old data engine (TDE):
vector-based engine inspired by MonetDB/X100

First step: Replace TDE as the backend of all Tableau Products

Goals:

- Deliver performance at scale
- Seamless transition for customers



TDE: The Gold Standard



Having a gold standard is great!

We just ran a lot (60k) of workbooks from Tableau Public and compared results.

Simple, measurable goal:
Get results to match for all and be fast,
then we're done 😊

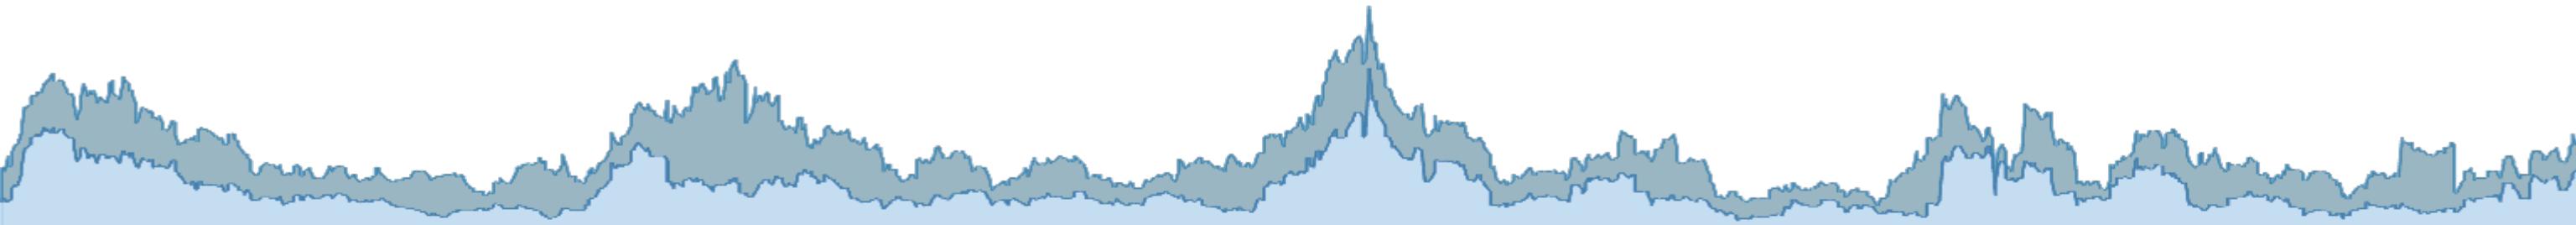
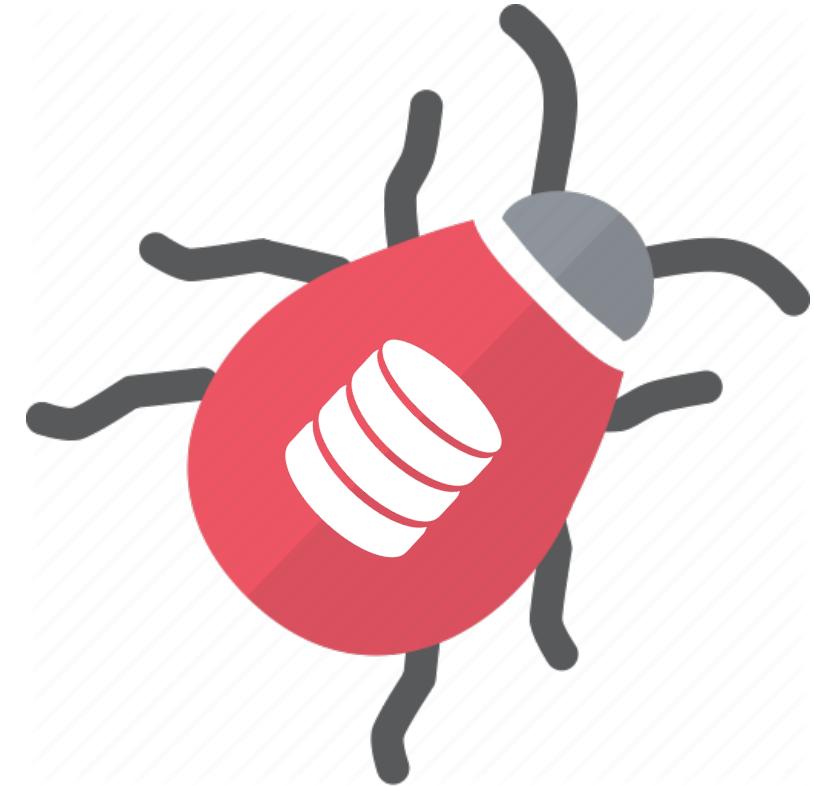
Challenge: Bug Compatibility?

Is it really worth to show the same result for all queries?

- Non-deterministic behavior (parallelization!)
- Bugs in TDE

Our attitude changed over time:

1. 100% same results at all cost, customers don't want their Viz to change!
This is non-negotiable!
2. Well, but what if it changes to be correct?
3. Who said that Visualizations can't change in the first place?
4. Let's do the right thing and fix things once and for all!



Compatibility Curiosity: String to Date Cast

Is “**5/7/2020**” April 7 or July 5?

TDE’s strategy: be aggressive finding a valid date. Sounds great!

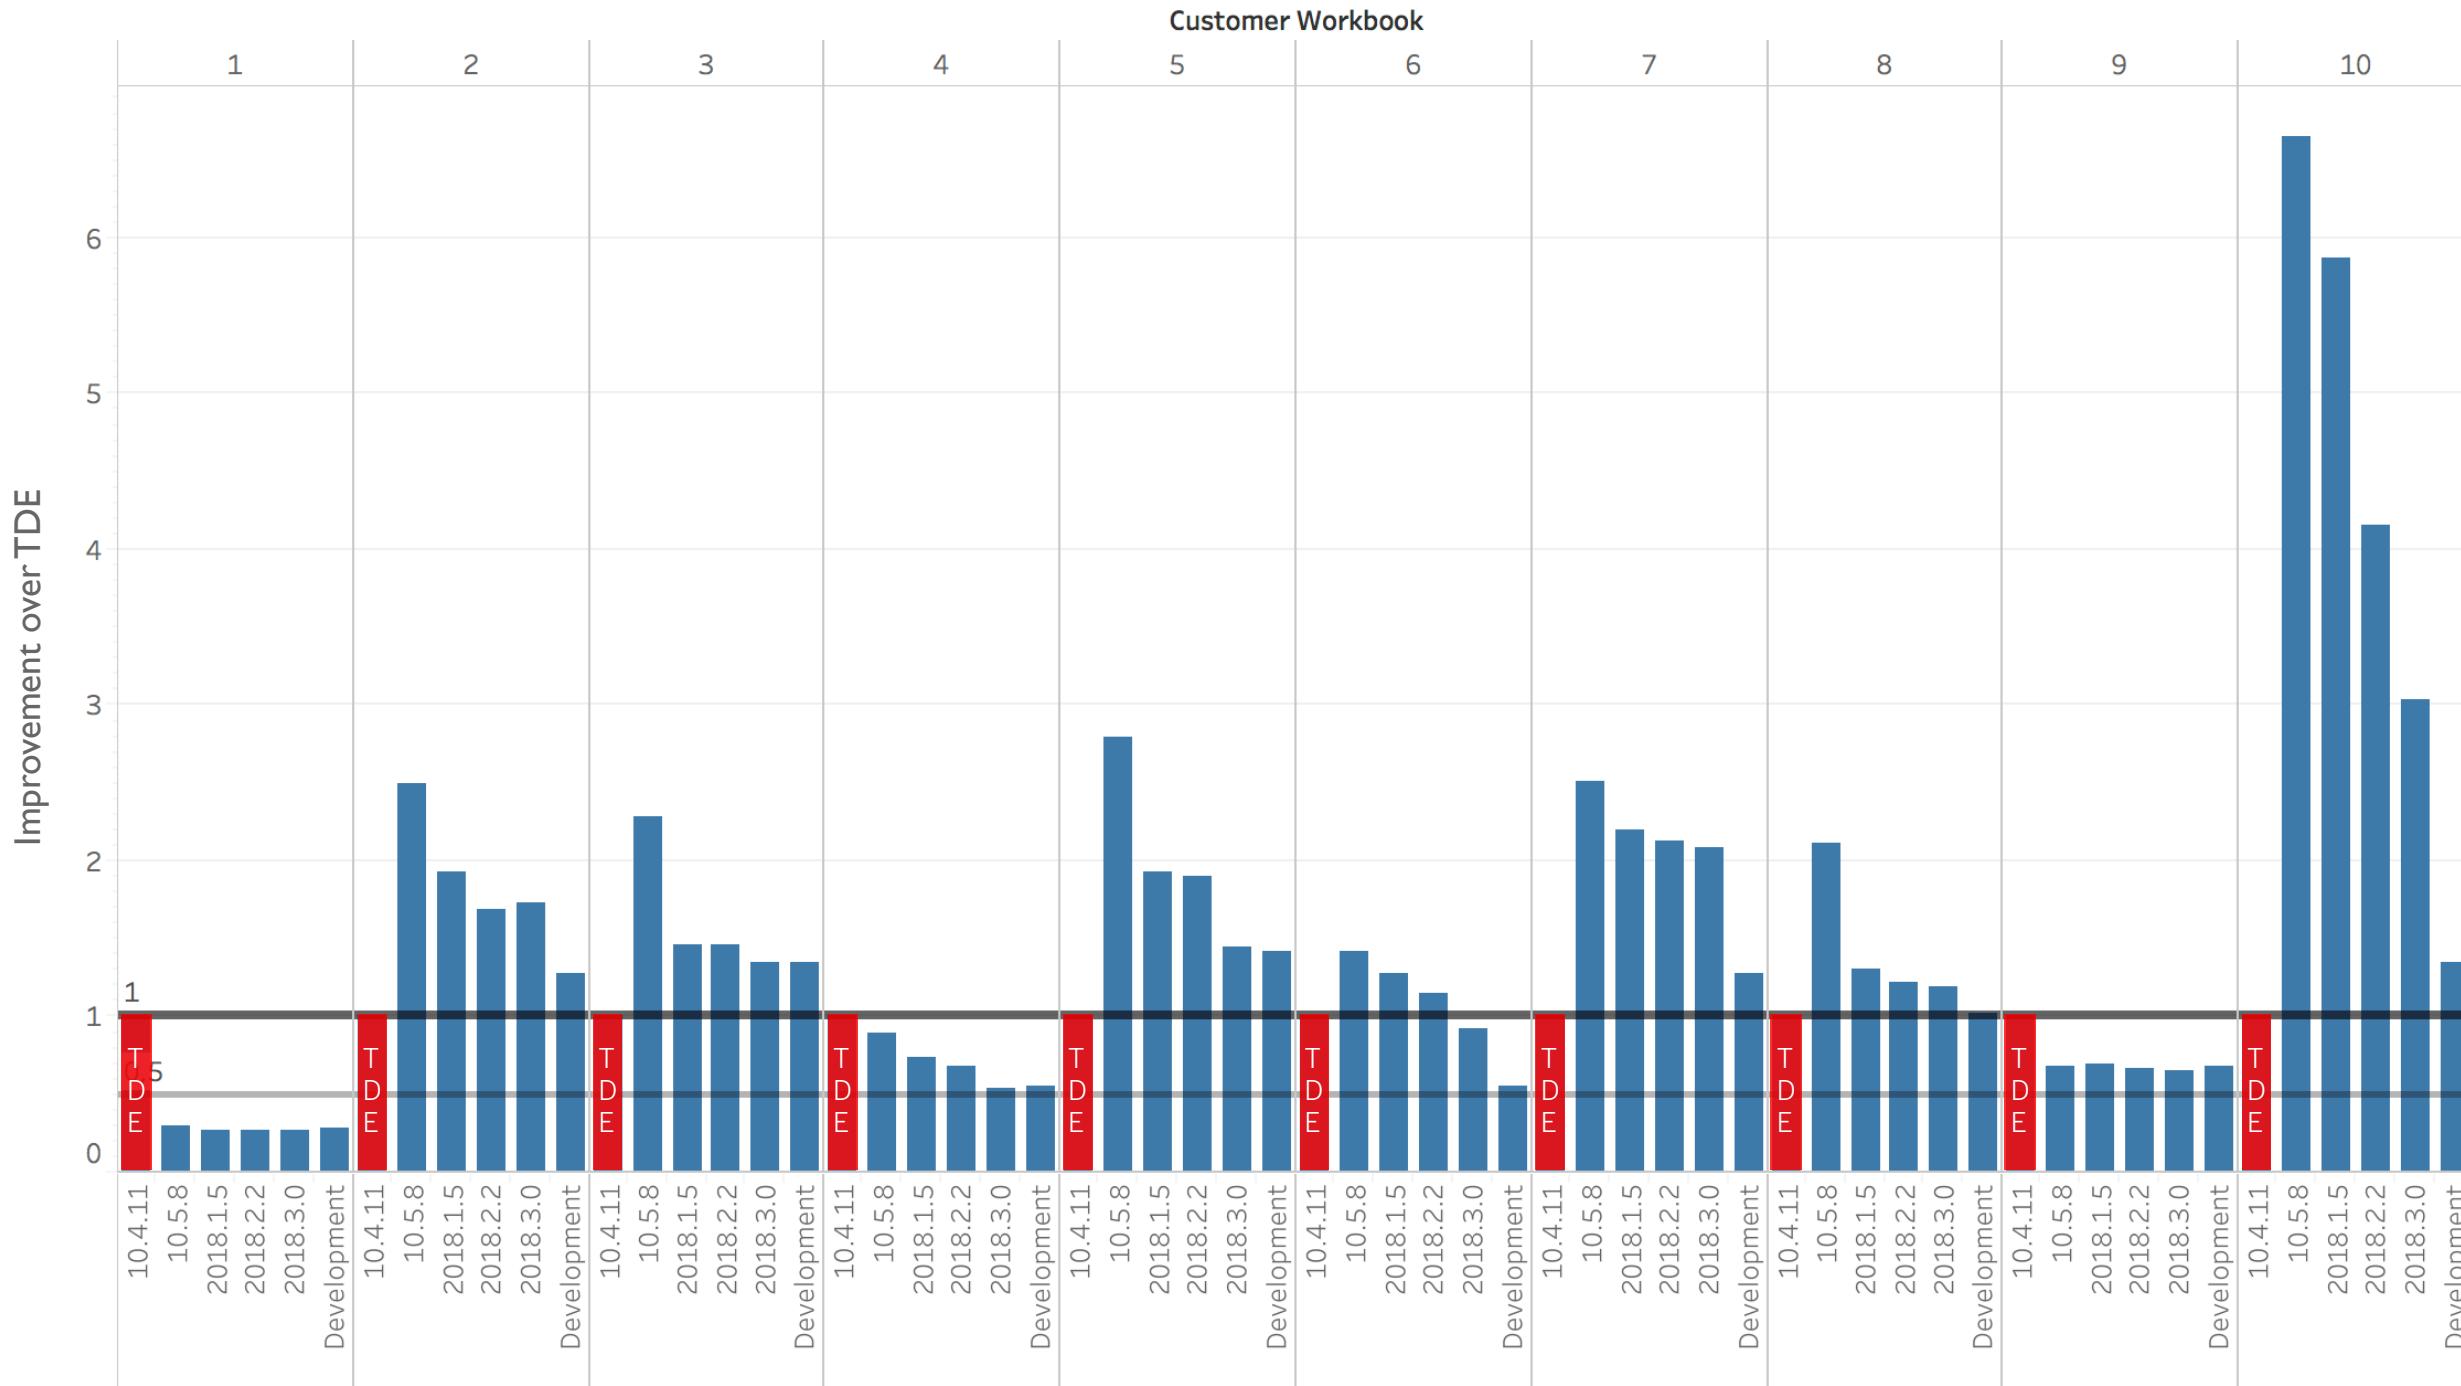
But horrible in the relational model!

Input	TDE	Hyper
5/7/2020	April 7, 2020	April 7, 2020
15/7/2020	July 15, 2020	NULL



Silent failure: Sales workbook: More sales in the first 12 days of each months!

Normalized query response times across major releases for reported regressions



Continuous performance improvements past launch

Testing Hyper



SQL Level Testing

We started with the SQLite test suite

- Added own test cases for features
- Added regression tests for defects
- Added fine grained expectations
(e.g., constant folding)

```
query N expectConstantResult
SELECT DATE '2001-09-28' + INTEGER '7'
-----
2001-10-05

# DATE + INTEGER -> DATE with overflow
statement error 22003
SELECT DATE '4713-01-01 BC' - 1
```

How to execute the tests?

- First: Own Hyper front end that parses the file and executes the queries:
 - Problem: Server / protocol code not tested
- Second: Client (based on libpq) that parses the file and sends queries to a server
 - Problem: Harder to debug, test driver is not same process

Testing even more with SQL Level Tests

Use EXPLAIN statement to test optimizer

Introduce function to scan the own log to test for log messages

- Introduce trace settings that allow printing specific internals to the log

Special test functions with side effects to test further internals

- E.g., a function that allocates thread-local memory
 - **SELECT suicide()**

```
# Deduplication for simple domain queries
query S
EXPLAIN SELECT * FROM
(SELECT a FROM t GROUP BY 1) t1,
(SELECT a FROM t GROUP BY 1) t2
-----
      executiontarget(1)
          join(2)
          bnl
explicitscan(3) explicitscan(6)->(4)
groupby(4)
tablescan(5)
t
```

SQL vs. C++ Unit Tests



SQL

Easy to write, usually very succinct

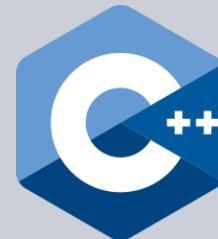
No recompilation needed

In vivo: Can run code in specific query contexts

Good test failure reporting

Can update expected test results automatically

Resembles customer usage of the system



C++

Only test the code in question,
not the whole SQL layer

Runs faster

Big controversy!

Beyond SQL Testing

SQLite tests are great, but they **can't simulate load** from multiple connections

Solution: Loadtest DSL

- open connection
- Embed SQLite test statement
- Execute code blocks in parallel
- Loops

```
exec CREATE DATABASE mytestdb;

connection mytestdb user=bob {
repeat 100 {
    exec CREATE TABLE foo AS
        (SELECT x FROM
         generate_series(1,1e6) x);
parallel 10 {
    exec UPDATE foo SET x = x+10
        WHERE x % 2 = 0;
} and 2 {
    test {
        query N
        SELECT SUM(x) FROM foo
        WHERE x % 2 = 1;
        -----
        1234567
    }
}
}
```

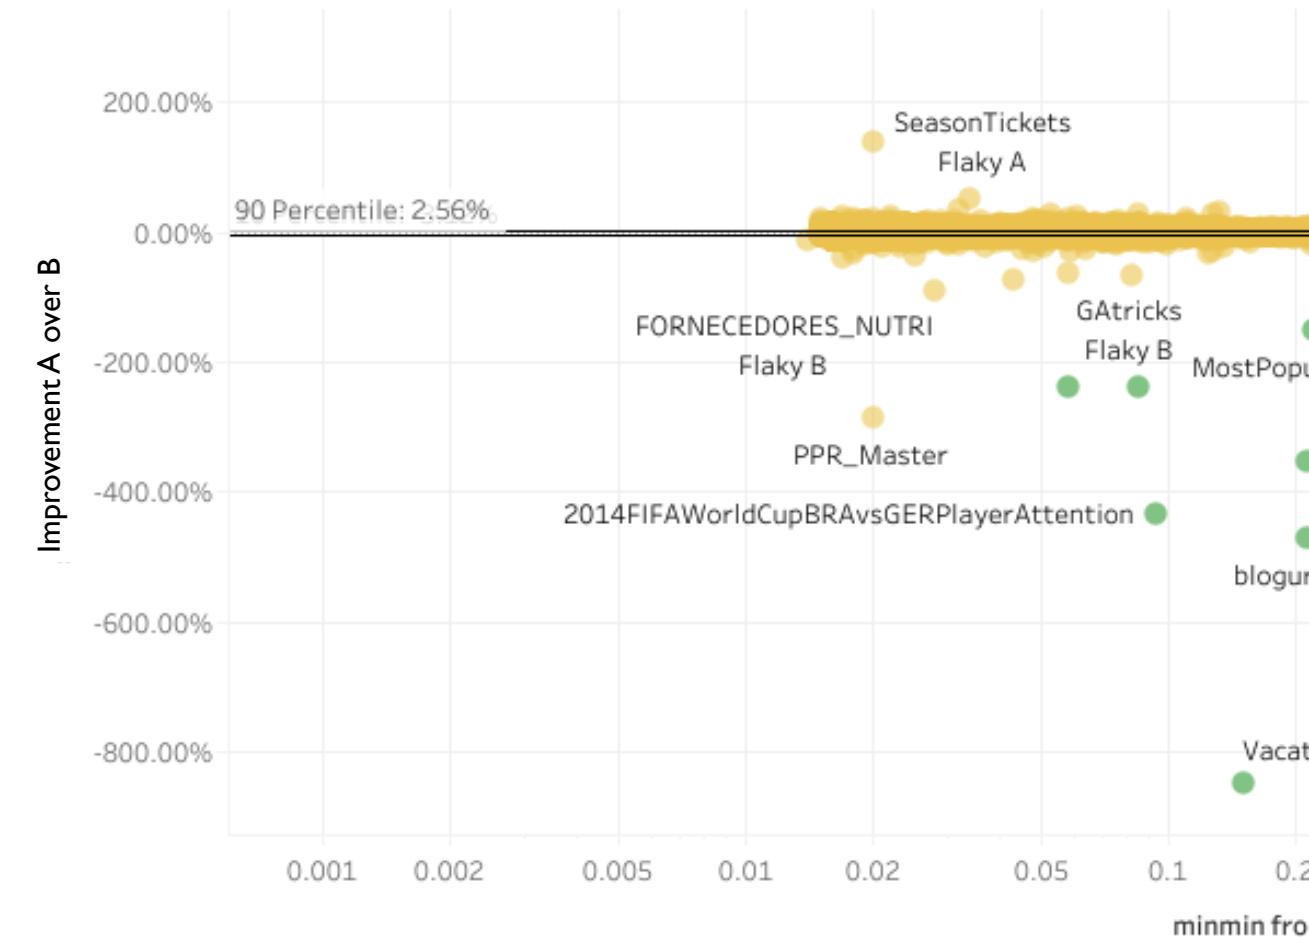
A/B Testing: MaxPerf / QueryRunner

1. Goal: Check for compatibility with TDE

- I. Correctness and perf
- II. A/B Test on the 60k Public Workbooks

2. How it works

- I. Starts up Hyper
- II. Loads a workbook
(thus sending queries to Hyper)
- III. Checks number of marks
/ mark values and records times
- IV. A/B test between old and new branch

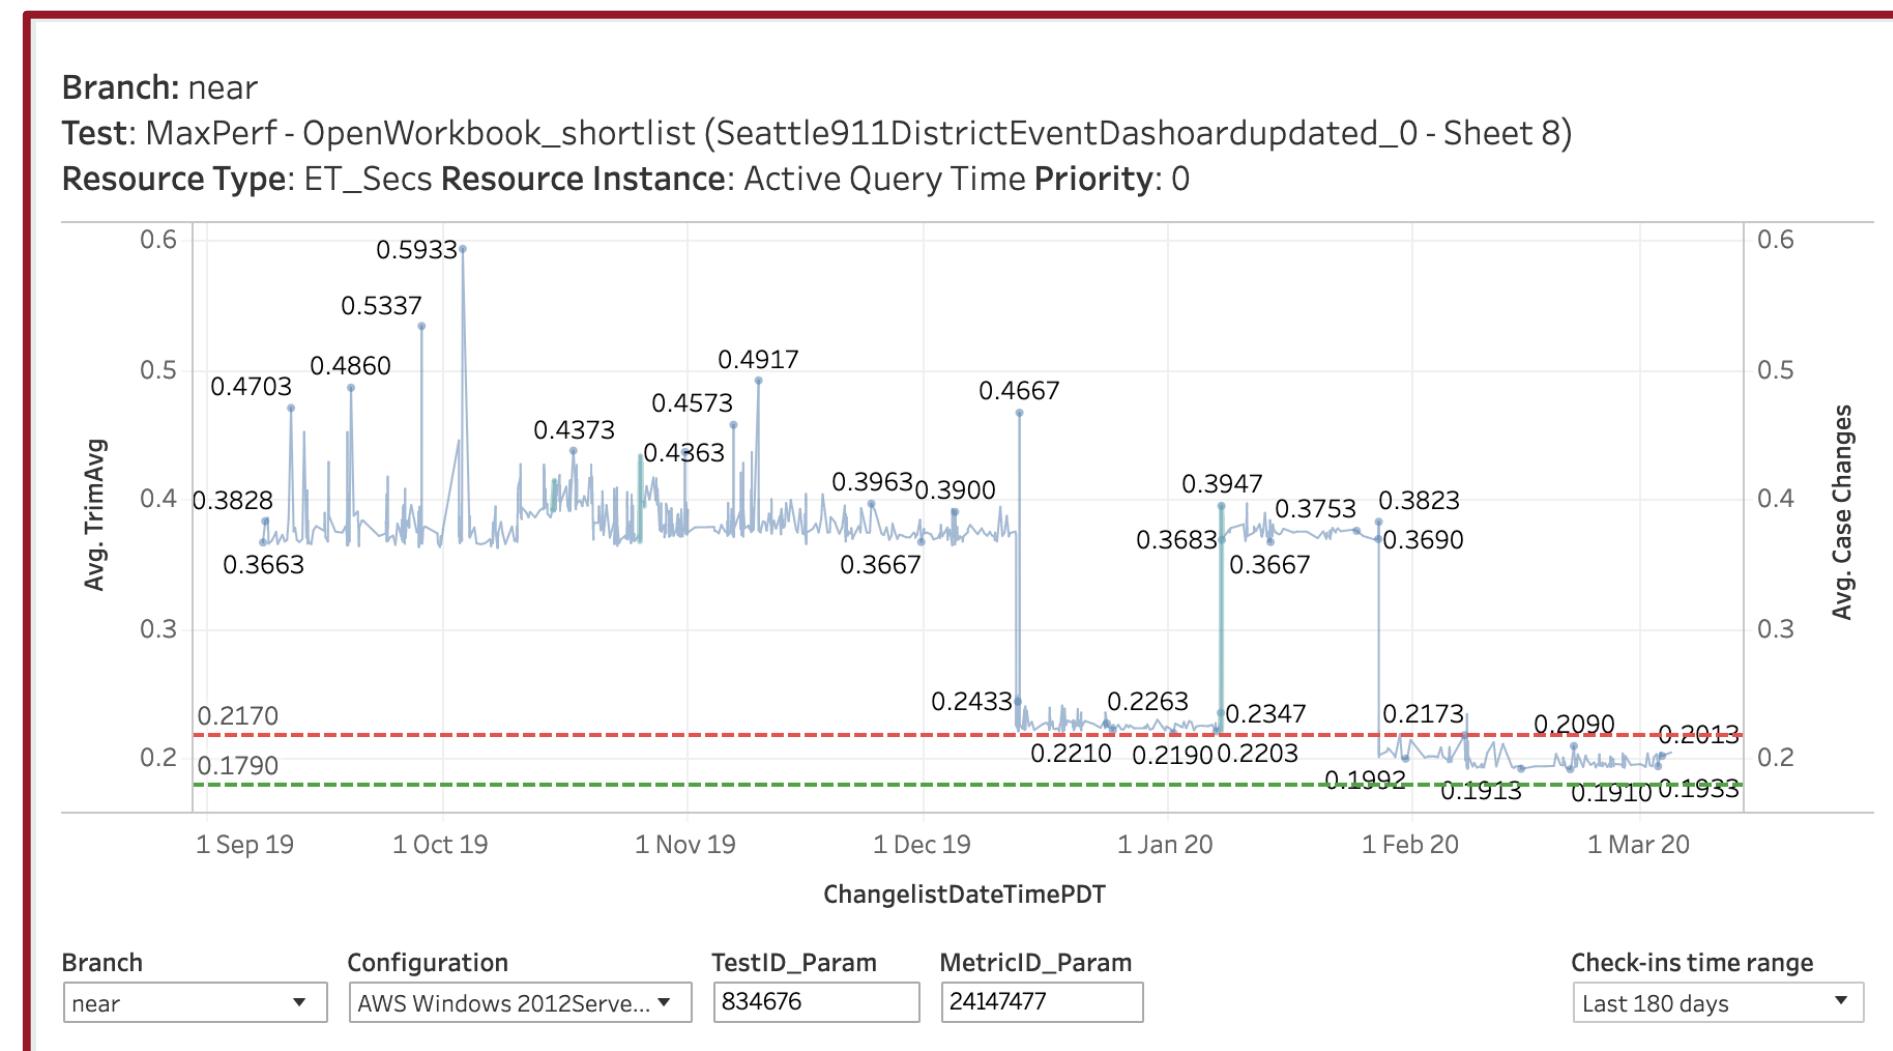


Automatic Regression Testing

DEFECT 1023800

1023800 Perf: 21 cases regressed by up to 356.67ms (up to 32%) starting from CL 1982374, near.19.1025.1835 🦆≠💀

1. Measure perf on every commit
2. If perf regresses, file a defect
3. Make sure it's not just noise
4. If perf improves, make it the new expectation



Fuzzing

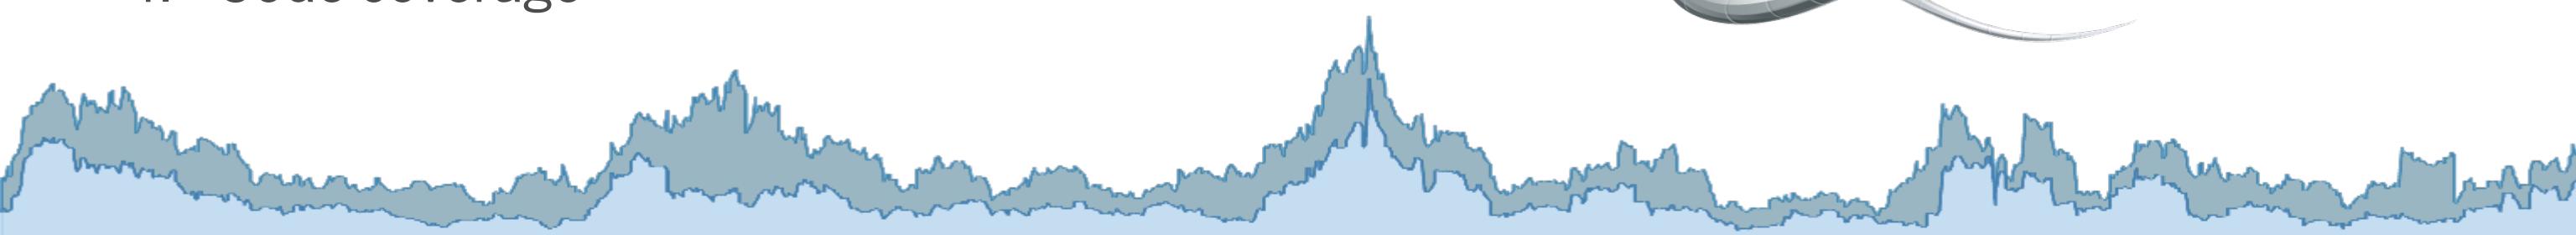
1. Use a fuzzer (e.g., AFL)
2. Feed it the SQL grammar
3. Let it run for a long time

american fuzzy lop 1.86b (test)	
process timing run time : 0 days, 0 hrs, 0 min, 2 sec last new path : none seen yet last uniq crash : 0 days, 0 hrs, 0 min, 2 sec last uniq hang : none seen yet	overall results cycles done : 0 total paths : 1 uniq crashes : 1 uniq hangs : 0
cycle progress now processing : 0 (0.00%) paths timed out : 0 (0.00%)	map coverage map density : 2 (0.00%) count coverage : 1.00 bits/tuple
stage progress now trying : havoc stage execs : 1464/5000 (29.28%) total execs : 1697 exec speed : 626.5/sec	findings in depth favored paths : 1 (100.00%) new edges on : 1 (100.00%) total crashes : 39 (1 unique) total hangs : 0 (0 unique)
fuzzing strategy yields bit flips : 0/16, 1/15, 0/13 byte flips : 0/2, 0/1, 0/0 arithmetics : 0/112, 0/25, 0/0 known ints : 0/10, 0/28, 0/0 dictionary : 0/0, 0/0, 0/0 havoc : 0/0, 0/0 trim : n/a, 0.00%	path geometry levels : 1 pending : 1 pend fav : 1 own finds : 0 imported : n/a variable : 0
[cpu: 92%]	

Found several vulnerabilities and defects with fuzzing!

Static Code Analysis

1. Enable all compiler warnings and make them errors
 - Wall -Wextra -Werror -Woverloaded-virtual -Wunreachable-code-return...
2. Keep code clean: clang-format, clang-tidy
3. Clang static analyzers
 - Address sanitizer
 - Thread sanitizer and Memory sanitizer
 - Third party libraries must also be re-built!
 - Undefined behavior sanitizer
 - ...
4. Code coverage

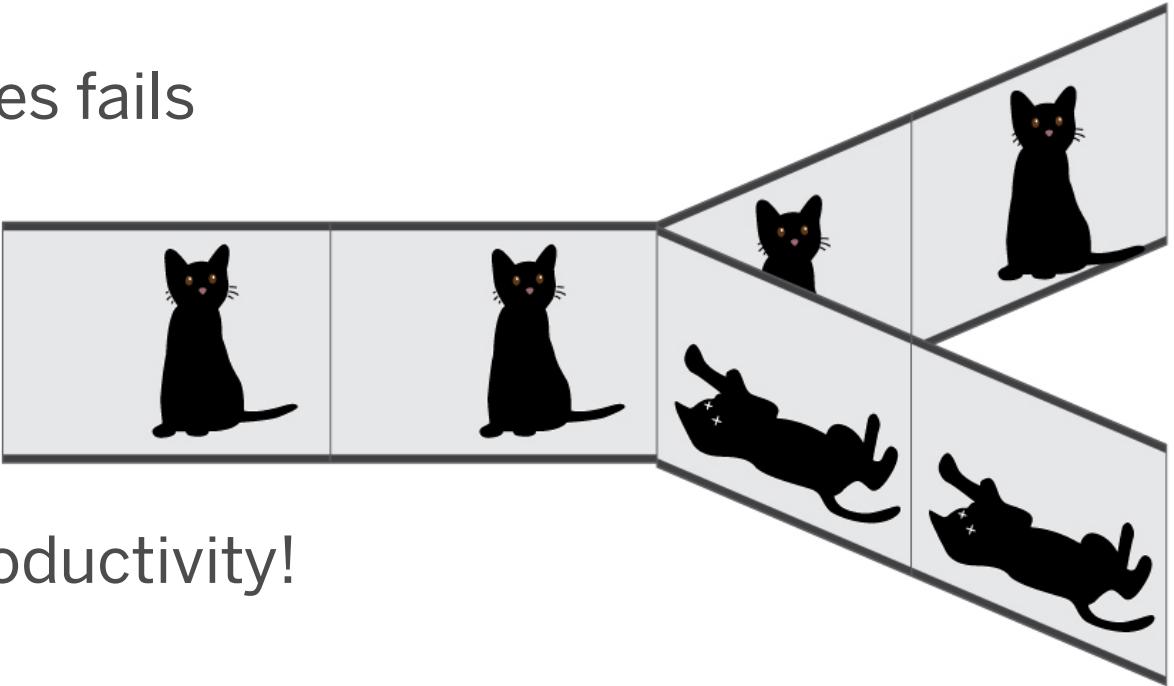


Flaky Tests: The root of all (test) evil

Flaky test: A test that sometimes passes and sometimes fails

Worst cases: It succeeds 99,9% of the time

- Don't let flaky tests build up!
- Don't get into the habit of muting flaky tests!
- Treat them as a high priority defects! They kill dev productivity!



Root causes:

- Real defect
- Bad test; usually dependent on timing or other

Lessons Learned

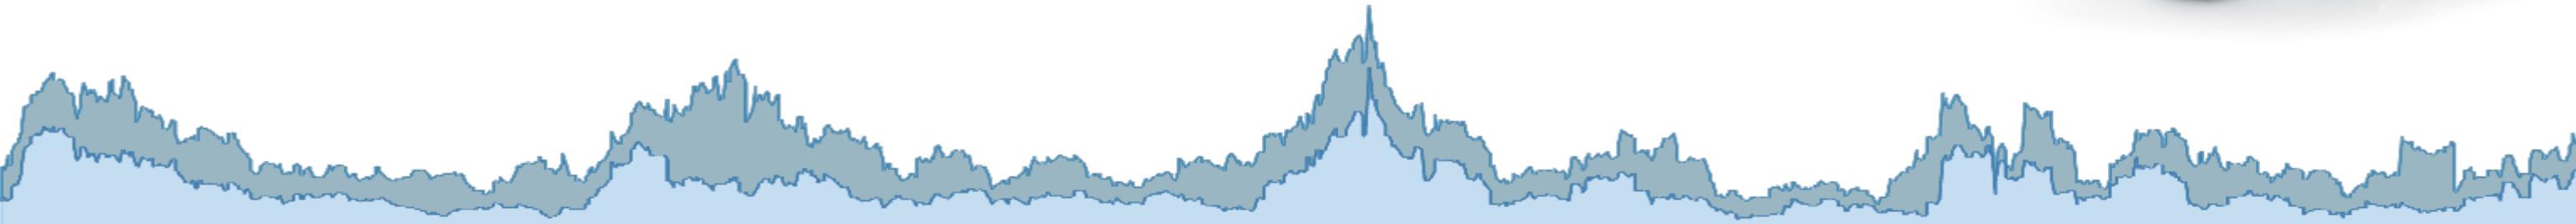
A classroom setting where several students are raised their hands, likely to answer a question. The students are seen from behind or side, wearing various clothing like a plaid shirt and a white sweater. The background is blurred, showing other students and classroom elements.

A classroom scene with several students raising their hands, likely to answer a question. The students are seen from behind or side, wearing various clothing like a plaid shirt and a white sweater. The background is blurred, showing other students and classroom elements.

Benchmark Responsibly!

When benchmarking a system ...

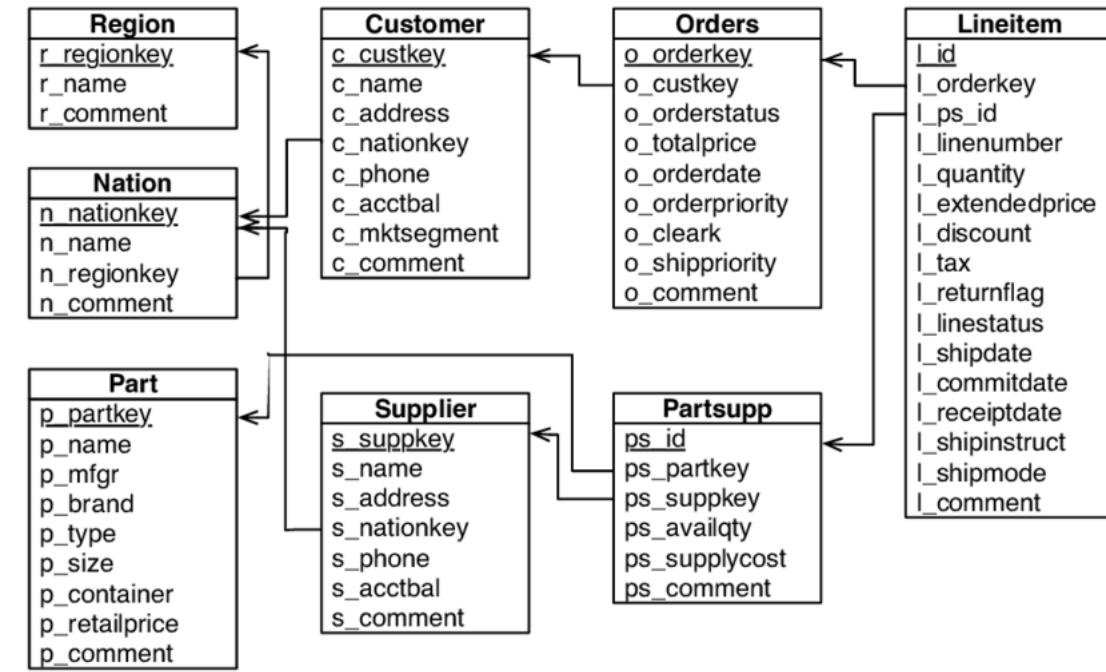
- Contact the vendor
- Report bugs
- Share your benchmark ahead of time, if possible
 - Allows vendors to give feedback, double check the validity
 - Good chance to increase the quality of your benchmark
- Ask the vendor how to configure the system
 - Don't let misconfigured benchmarks impact the credibility of your hard work
- Shout-out to Andrew C. from Brown University



There is more than TPC-X

Generated workloads are real

- I. More and more tools generate queries
- II. Queries are way more complex than hand-written TPC-X queries
- III. Plenty of interesting (and novel!) problems lurking in other workloads



Public BI Benchmark from CWI

- https://github.com/cwida/public_bi_benchmark
- Ghita, Tomé, Boncz, **White-box Compression: Learning and Exploiting Compact Table Representations**, CIDR'20
- Based on Tableau Public data

Build your system for production

Build a system; not a throw-away prototype

- Much more rewarding, longer lasting sense of achievement
- Easier to build upon previous work
- Results closer to reality (micro benchmarks leave out crucial parts)

Follow standards (e.g., SQL, PostgreSQL)

- Easy test adaption, easier benchmarking

Build that system **as if it was for production**

- Test driven development is great;
defects could make your results invalid!
- Easier to adapt a stable system to the next benchmark
- There is a real chance to get your system into production in the end
- OS compatibility (Some perf hacks are highly non-portable!)



Conclusion



Conclusion

Academic Projects can make it into production

- Design to become a product (it's fun!)
- Replacing an old system: great gold standard, hard to beat in all cases

TPC-X is not everything!

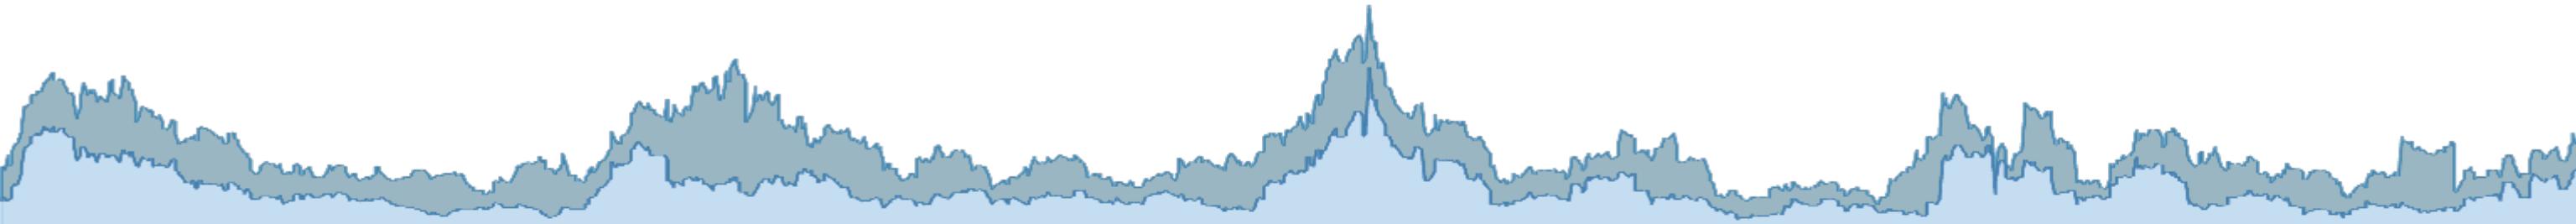
- Experiment with more diverse workloads, real world queries are complex

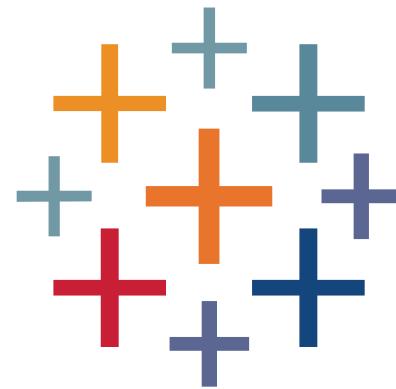
Various layers of testing required

- Don't forget stress testing, don't regress, and avoid flaky tests

Benchmark responsibly

Try out Hyper API! tabsoft.co/hyperapi





+ a b l e a u®